

# CMPT 300 D100

## Operating System I

### 1.2 OS Introduction

**Dr. Hazra Imran**

**Spring 2022**

# Learning goals

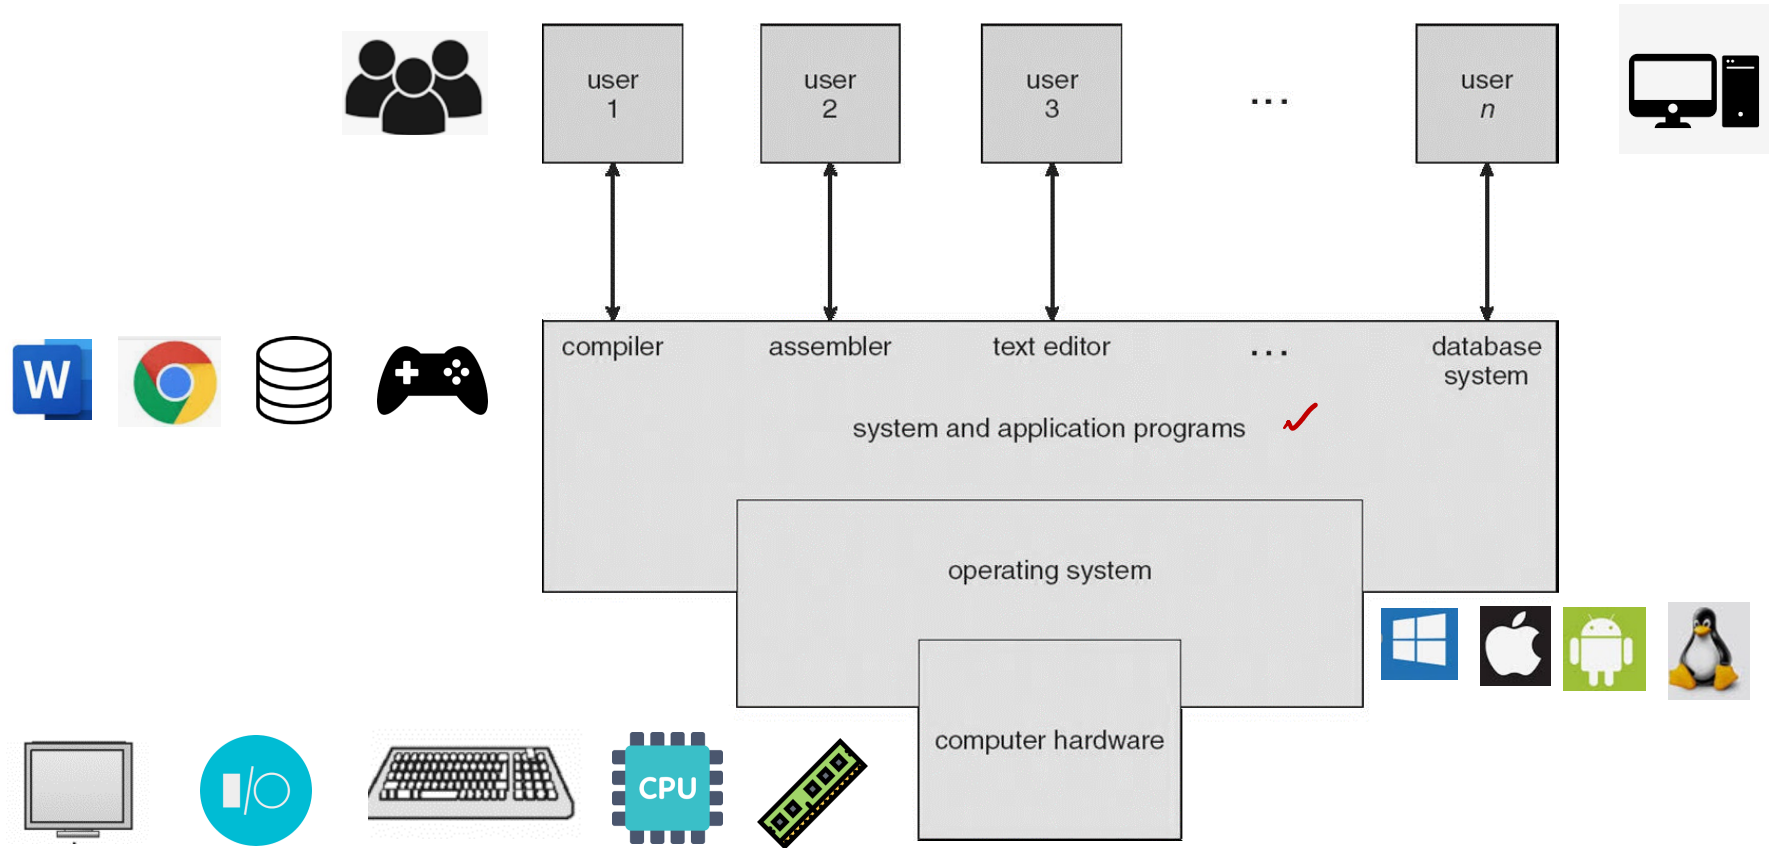
- Describe the general organization of a computer system
- Describe the role of interrupts.

# Computer System Structure

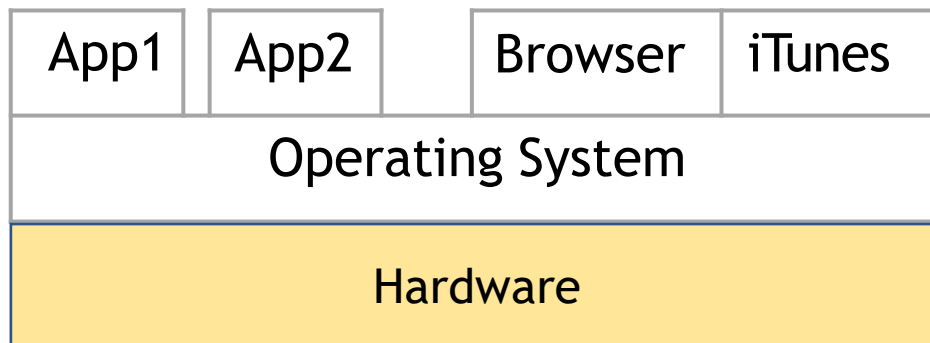
*processes*

Computer system can be divided into four components:

*job*



# Why Start With Hardware?



- OS functionality depends upon the architectural features of H/W  
*enforce protection* *resource sharing*
- Architectural support can greatly simplify or complicate OS tasks

# Role of Operating System

P1 ——— X  
P2 ——— X

- OS is a **resource allocator**
  - Manages all hardware resources
  - Decides between conflicting requests for fair resource use
- OS is a **control program**
  - Controls execution of programs to prevent errors and improper use of the computer
- OS provides **abstractions** ✓
  - Hides the details of the hardware
  - Provides an interface that allows a consistent experience for application programs and users



GUI  
CLI

# Example of abstractions

- Allows user programs to deal with simpler, high-level concepts
- Hides complex and unreliable hardware
- Provides illusions like “only one application running” or “infinite memory available”

Runny

→ P<sub>1</sub> —

P<sub>2</sub> —

P<sub>3</sub> —

# What if we read directly from the files?

- Where is this file on disk? Which track and sectors?
- Code needs to change on a different system

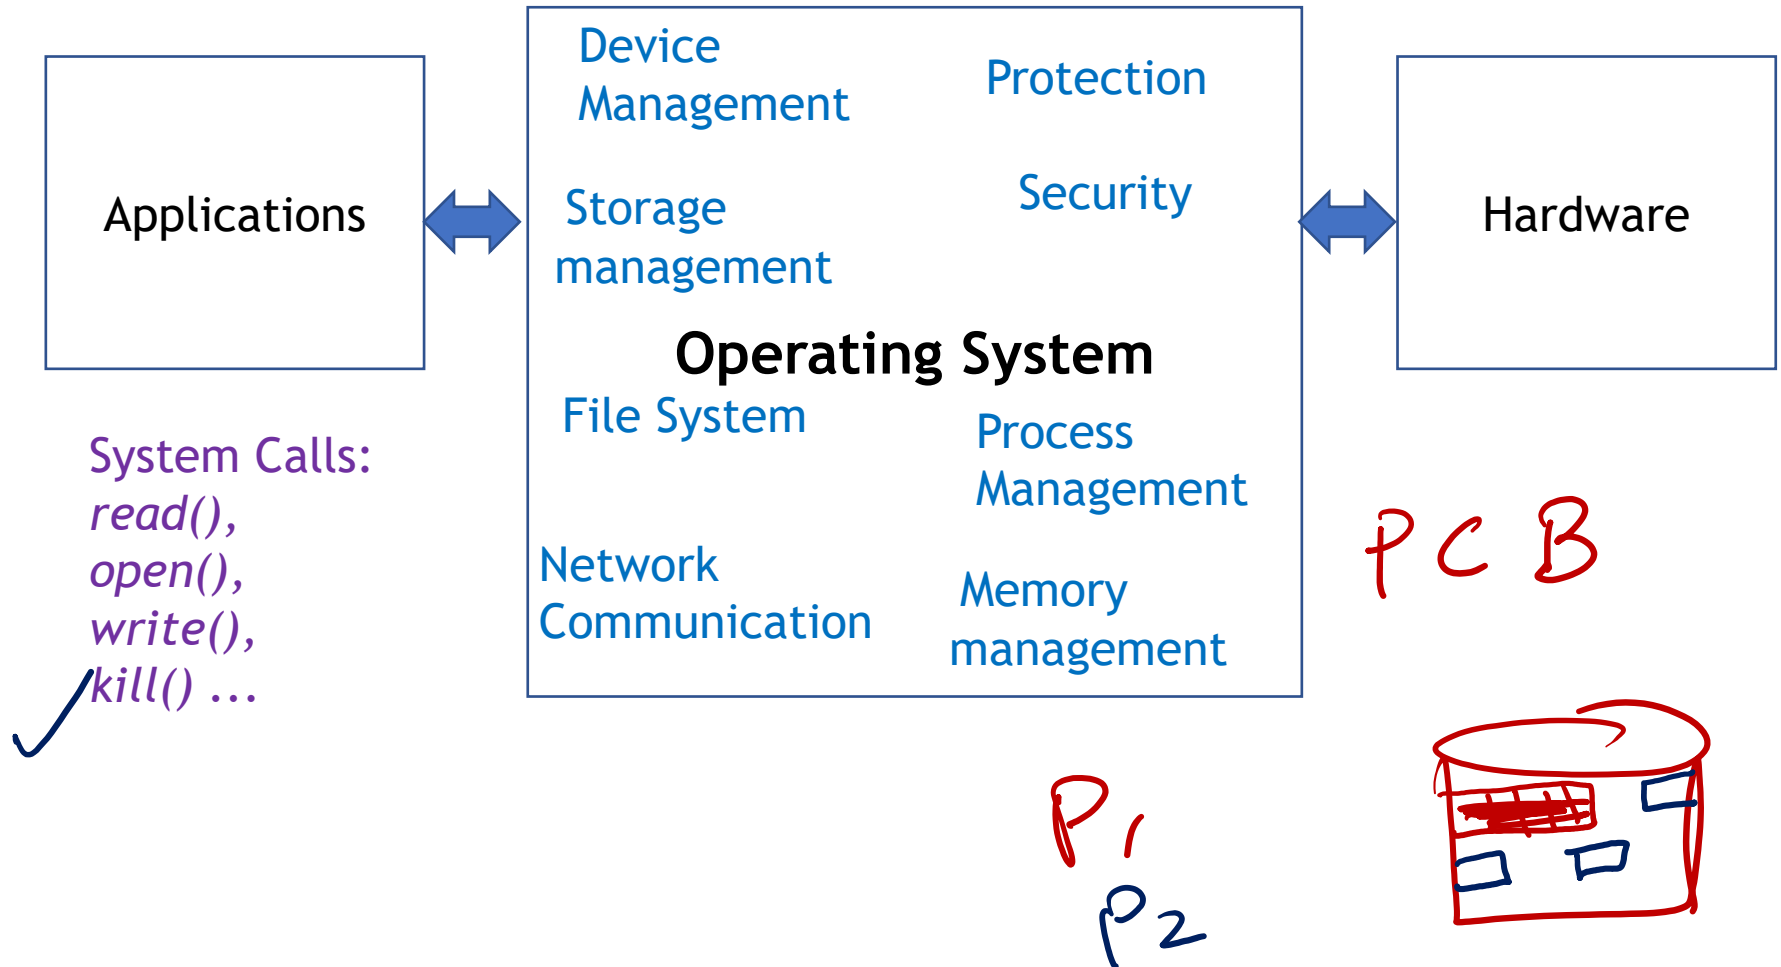


Abstraction ....

```
file = open ("test.txt")  
...  
close (file);
```

*detail / commands.*

# Example of abstraction





# Try out Activity - 1

load  
store add

Hardware component	Hardware capabilities	User needs
CPU	Machine instructions - perform operations on contents of registers and memory locations.	User thinks in term of <ul style="list-style-type: none"><li>• arrays high level</li><li>• list ds</li></ul>
Main memory	Physical memory is a linear sequence of addressable bytes or words that hold programs and data.	<ul style="list-style-type: none"><li>• var - heap</li><li>• executables lib</li></ul>
Secondary storage	Are multi-dimensional structures, which require complex sequences of low-level operations to store and access data organized in discrete blocks.	programs to manipulate / access data
I/O Devices	are operated by reading and writing registers of the device controllers.	Interfaces

# Instruction Cycle

The CPU repetitively performs the instruction cycle:

## Fetch

- The program counter (PC) holds the address in memory of the next instruction to execute
- The address from memory is fetched and stored in the IR
- The PC is incremented to fetch the next instruction (unless told otherwise)


## Decode

- The CPU determines what instruction is in the IR

10011010100000  
stop  
word


## Execute

- Circuitry interprets the opcode and executes the instruction
- ✓ Moving data, performing an operation in the ALU, etc.
- May need to fetch operands from memory or store data back to memory




# Single-core processor vs Multiple-core processor

Single-Core Processor



Multi-Core Processor





- ☹ High power consumption
- ☹ Heat generation
- ☹ 75% of CPU time is used waiting for memory access results.
- ☹ Limited parallelism

# Massive parallelism offered by

## Hyper threading

- Each core may execute more than one thread
  - Hyper-threading
  - Share certain physical resources (e.g., caches, registers)



## Multi-core and multi-socket


- Socket to place a multicore processor
- Usually 4-32 cores per processor/socket




Process  
T1 { }  
T2  
T3

# Shared memory systems

## Centralized Shared Memory Architecture aka.. Symmetric Multi-Processors (SMP)



## Distributed Shared Memory Architecture aka.. Non-Uniform Memory Architectures (NUMA)




Memory local to a core is faster to access than non-local ~~memory~~

# Some useful system monitoring tools

- **htop** - highlights the processes and information using some color
- nmon - lets you monitor system resources and processes.
- CoreFreq - a CPU monitoring program
- <https://linuxide.com/monitoring-2/10-tools-monitor-cpu-performance-usage-linux-command-line/>

# Computer System Operation



- I/O devices and the CPU can execute concurrently
- Each device controller is in charge of a particular device type
- Each device controller has a local buffer
- CPU moves data from/to main memory to/from local buffers
- I/O is from the device to local buffer of controller
- Device controller informs CPU that it has finished its operation by causing an interrupt

ISR



# Interrupts

- An interrupt is a signal that causes the computer to alter its normal flow of instruction execution.
- Allows the processor to execute other instructions while an I/O operation is in progress
- Improves processing efficiency
- ✓ • A suspension of a process caused by an event external to that process and performed in such a way that the process can be resumed

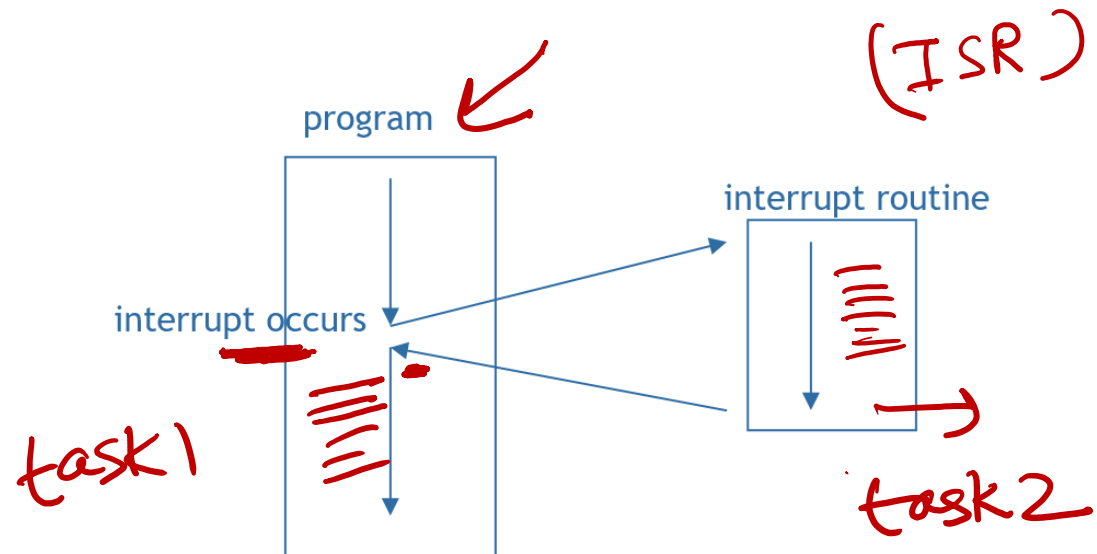
## Interrupt Handler

- A program that determines nature of the interrupt and performs whatever actions are needed
- Control is transferred to this program
- part of the operating system

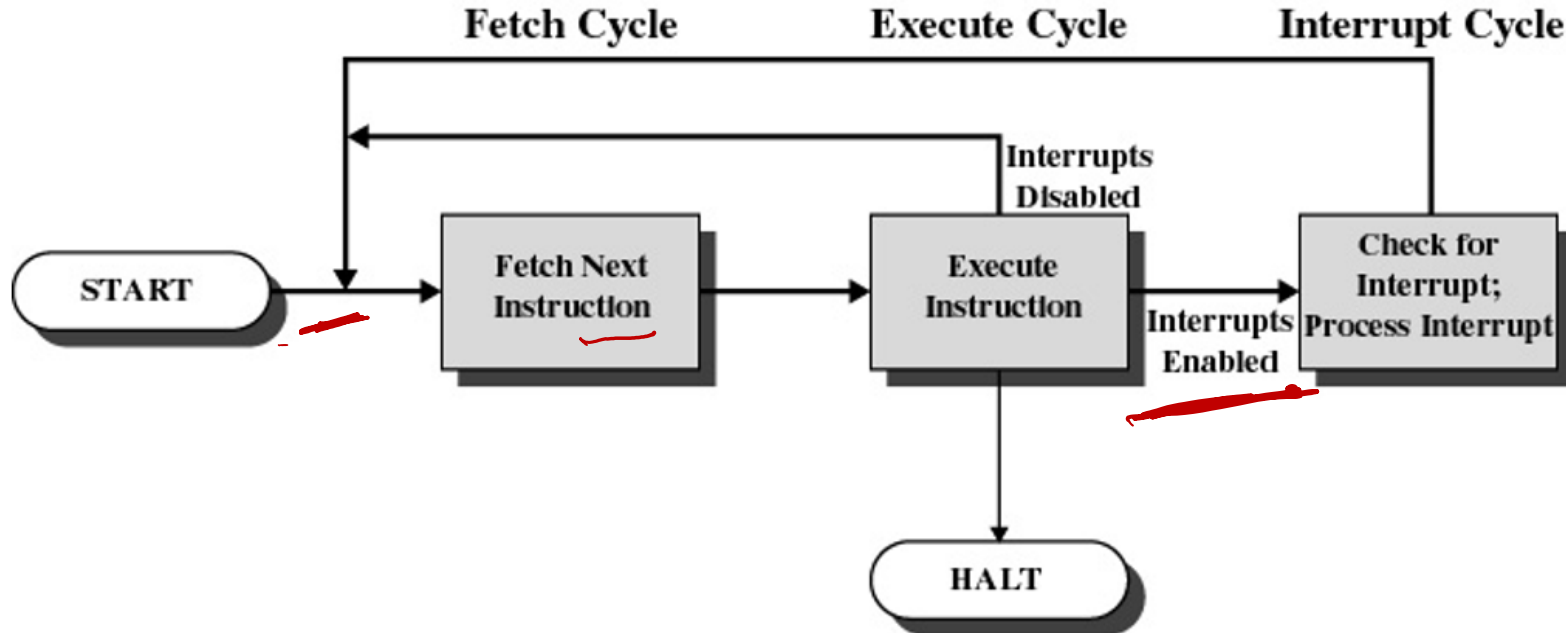


# Classes of Interrupts

- Program
  - arithmetic overflow
  - division by zero
  - execute illegal instruction
  - reference outside user's memory space
- Timer ✓
- I/O
- Hardware failure



# Interrupt Cycle



- Processor checks for interrupts
- If no interrupts fetch the next instruction for the current program
- If an interrupt is pending, suspend execution of the current program, and execute the **interrupt handler**

# Interrupt Handling

- The operating system preserves the state of the CPU by storing registers and the program counter
- Determines which type of interrupt has occurred:
  - **polling**
  - **vectored** interrupt system
- Separate segments of code determine what action should be taken for each type of interrupt


*“ Polling is like opening the door every few seconds to see if guests have arrived...”*



*Interrupts are like waiting for the door bell to ring.*




# Interrupt Timeline



# How to Find the “Corresponding ISR”?

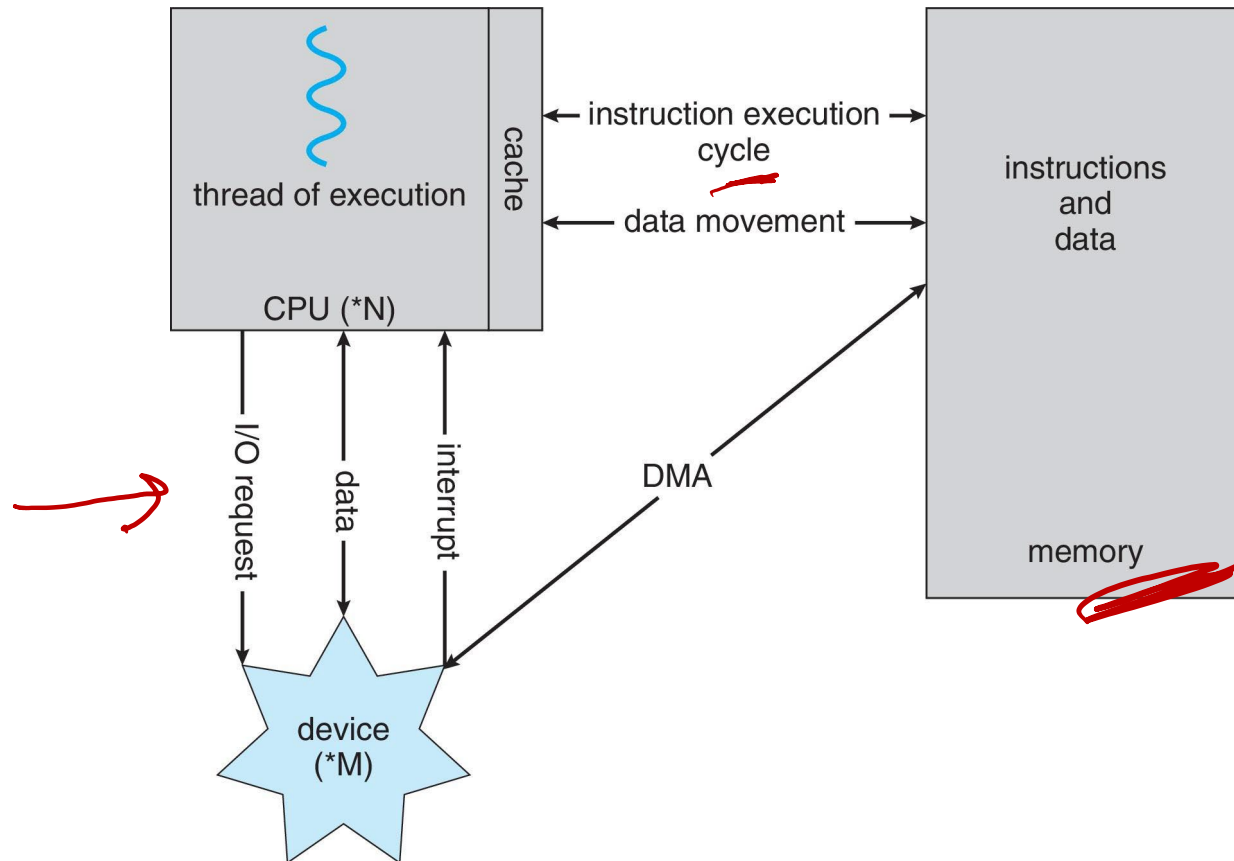
- Interrupts must be handled very quickly
- The system can support multiple interrupts
- Need multiple different ISRs
- Interrupt vector
- Table indexed by device number
- Entry contains address to the ISR



# I/O Structure

- After I/O starts, control returns to user program only upon I/O completion
  - Wait instruction idles the CPU until the next interrupt
  - Wait loop (contention for memory access)
  - At most one I/O request is outstanding at a time, no simultaneous I/O processing
- After I/O starts, control returns to user program without waiting for I/O completion
  - **System call** - request to the OS to allow user to wait for I/O completion
  - **Device-status table** contains entry for each I/O device indicating its type, address, and state
  - OS indexes into I/O device table to determine device status and to modify table entry to include interrupt

# Summary: How Modern Computer Works



# Try out activity

Program 1 and 2 run concurrently. Whenever a timeout interrupt occurs, the kernel switches control between the programs. Show the order of instruction execution, assuming program 1 is currently running.

Application 1	Application 2
...	
instruction i	instruction 0
(timeout interrupt)	...
instruction i+1	instruction j
...	(timeout interrupt)
instruction k	instruction j+1
(timeout interrupt)	...
instruction k+1	
...	

HHW