

CMPT 300

Operating System I

3.3 -Process 3
Chapter 3

Dr. Hazra Imran


Admin notes

- A2 available on canvas

PCB organization

Two ways

1. An array of structure




draw back ? wastage of memory

PCB organization

Two ways

2. An array of pointers to dynamically allocated PCBs




drawback
overhead.

PCB in Linux

Represented by the C structure `task_struct`

✓(`<include/linux/sched.h>`)

```
pid_t pid;           /* process identifier */
long state;          /* state of the process */
unsigned int time_slice /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm; /* address space of this process */
```



✓
current
(currently executing process)

current → state =
new-state

Process State Queues

- The OS maintains the PCBs of all processes in state queues.
- The OS maintains a queue for each of the states.
- PCBs of all processes in the same execution state are placed in the same queue.
- When the state of a process is changed, its PCB is unlinked from its current queue and moved to its new state queue.


Process Scheduling

- The goals are to:
 - Maximize CPU use ✓
 - Give processes the CPU time that they need
 - Process scheduler selects among available processes for next execution on CPU
- metrics, parameters*

Process Scheduler

Maintains **scheduling queues** of processes


- **Ready queue:** set of all processes residing in main memory, ready and waiting to execute




Process Scheduler

Maintains **scheduling queues** of processes

- **Ready queue:** set of all processes residing in main memory, ready and waiting to execute



- **Wait queues**



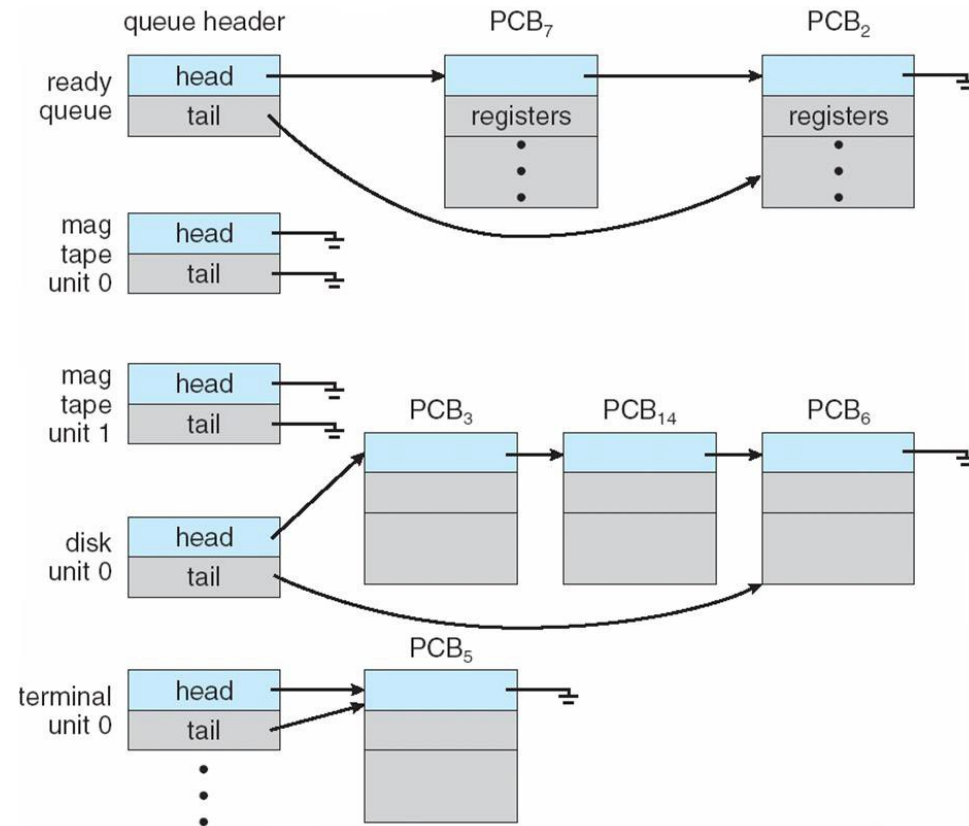
- Processes migrate among the various queues, depending (in part) on their state

Ready Queue And Various I/O Device Queues

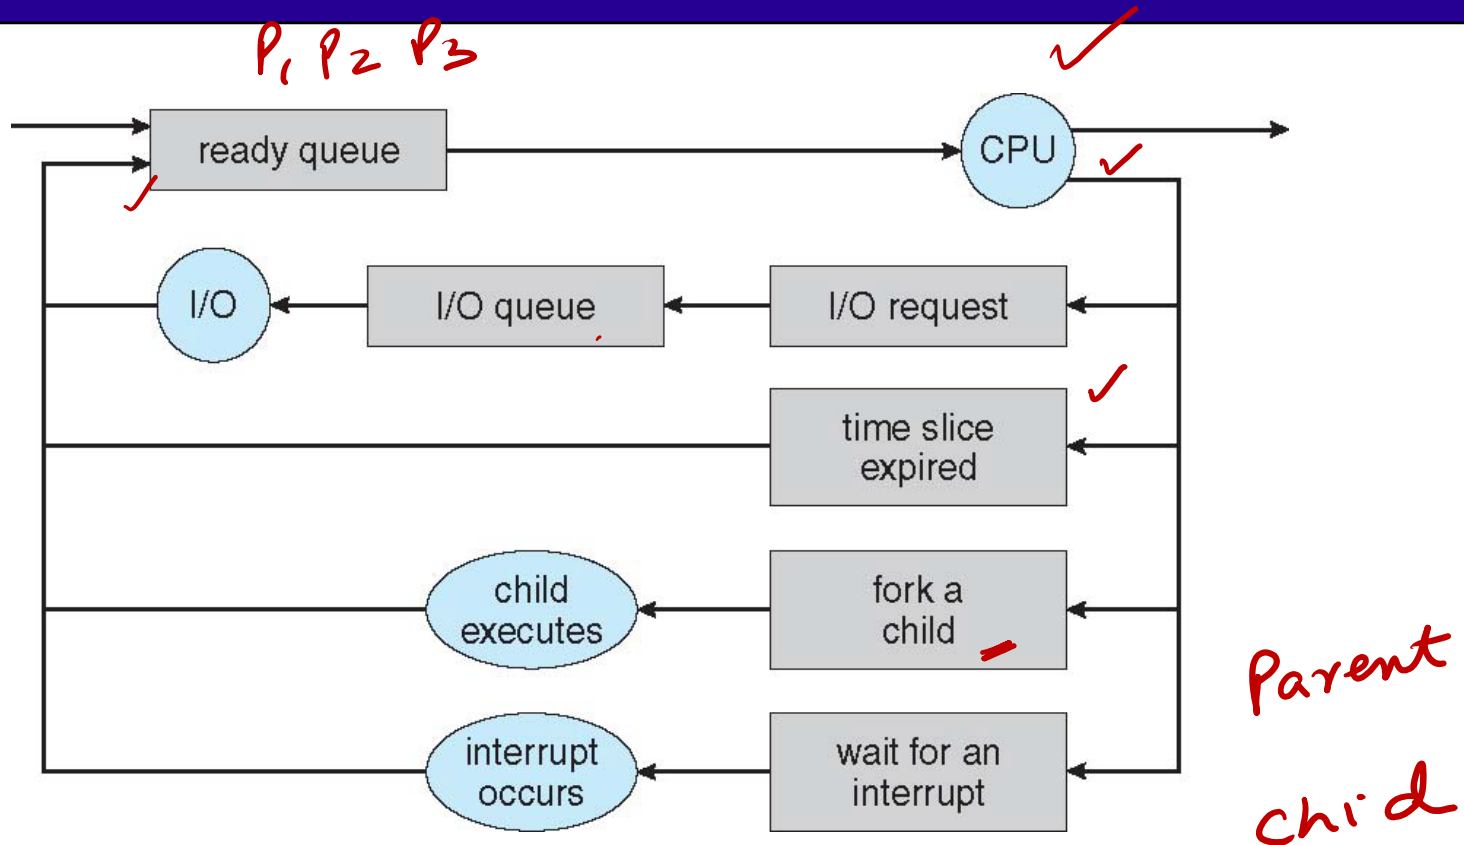
Processes are maintained in queues

Ready queue - set of all processes residing in main memory, ready and waiting to execute

Device queues - set of processes waiting for an I/O device



Scheduling Queues



Queueing diagram represents queues, resources, flows

I/O bound v/s
CPU bound processes

↓
spends more time doing computation
long CPU burst

Scheduler Components

Short-term scheduler (or CPU scheduler)

- selects which process should be executed next and allocates CPU
- is invoked frequently (milliseconds), so it must be fast

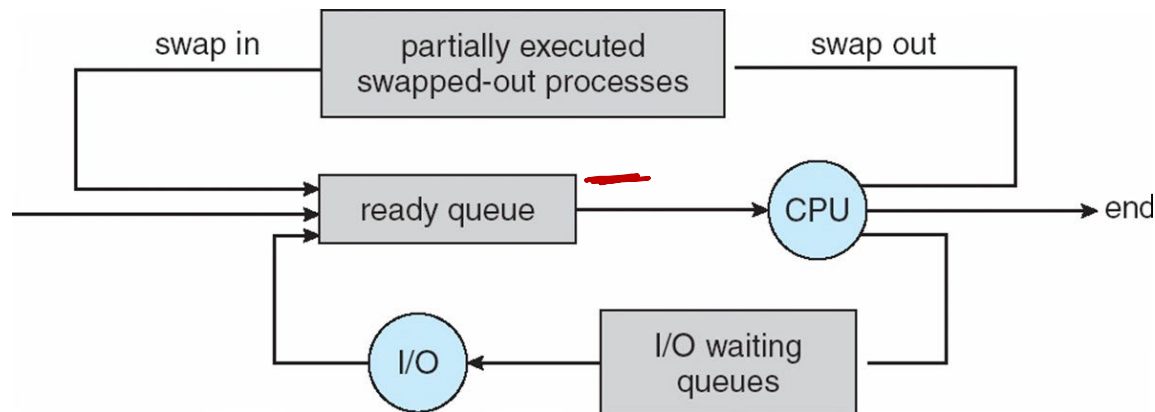
Long-term scheduler (or job scheduler)

- selects which processes should be brought into the ready queue
- is invoked infrequently (seconds, minutes), so it may respond slowly

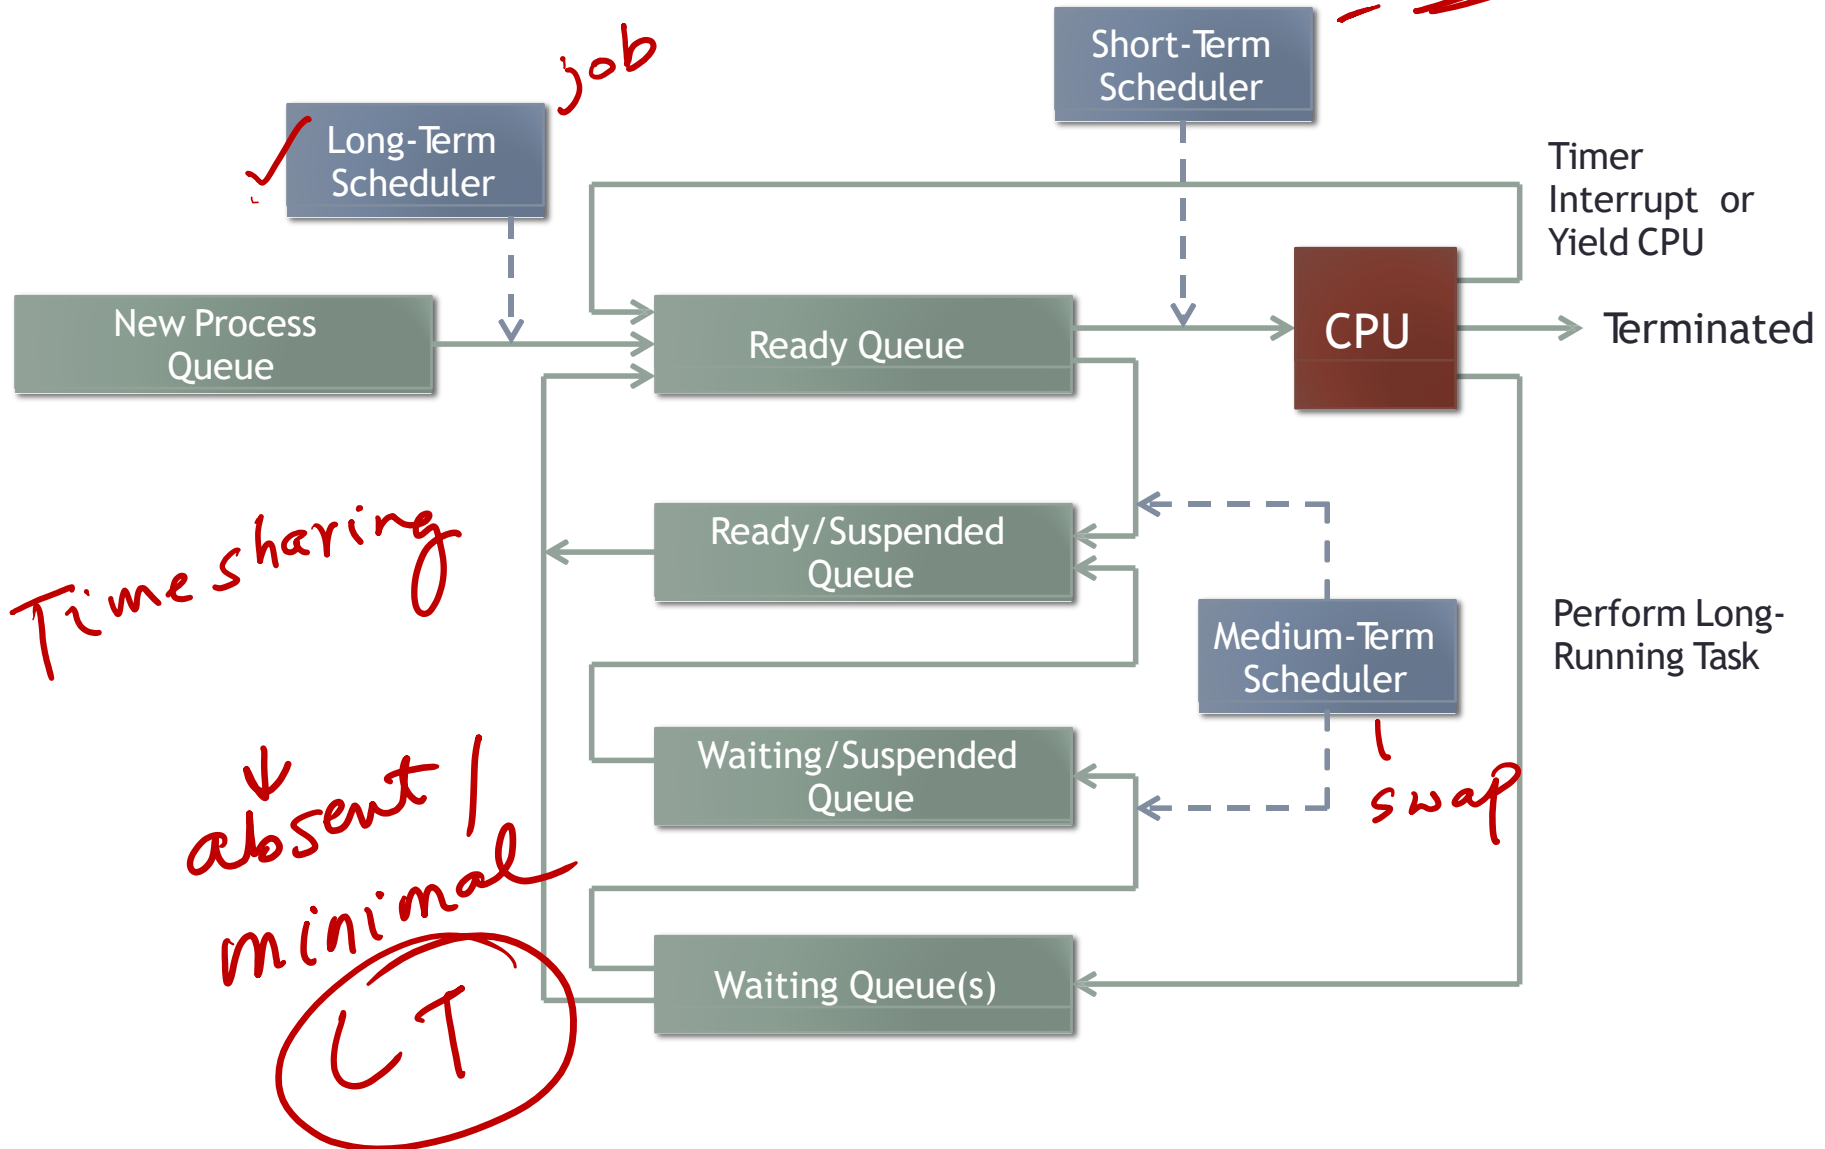
Medium Term Scheduling

Medium-term scheduler can be added if degree of multiple programming needs to decrease

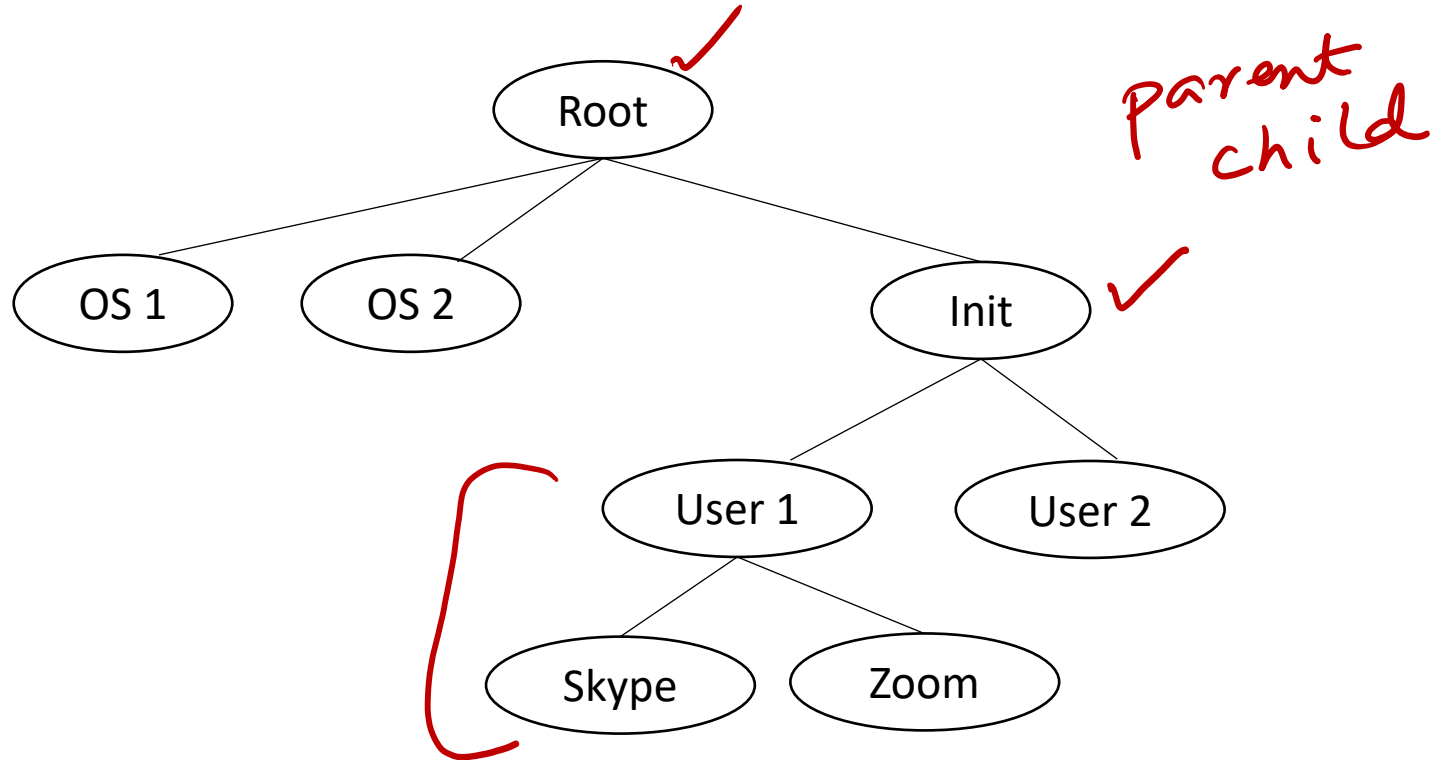
Remove process temporarily from memory, store on disk, bring back in from disk to continue execution: swapping



Process Scheduling



Process Creation Hierarchy



Operation on processes

✓ ps -el

✓
active

```
user1@user1-VirtualBox:~$ ps -el
F S      UID        PID      PPID    C  PRI   NI     ADDR  SZ  WCHAN    TTY          TIME CMD
4 S      0          1          0      0   2   80     0 - 41885  -    ?           ?      00:00:01 systemd
1 S      0          2          0      0   0   80     0 -      -    ?           ?      00:00:00 kthreadd
1 I      0          3          2      0   60  -20    -    0 -    ?           ?      00:00:00 rcu_gp
1 I      0          4          2      0   60  -20    -    0 -    ?           ?      00:00:00 rcu_par_g
1 I      0          5          2      0   80     0 -    0 -    ?           ?      00:00:00 kworker/0
1 I      0          6          2      0   60  -20    -    0 -    ?           ?      00:00:00 kworker/0
1 I      0          7          2      0   80     0 -    0 -    ?           ?      00:00:00 kworker/0
1 I      0          8          2      0   80     0 -    0 -    ?           ?      00:00:00 kworker/u
1 I      0          9          2      0   60  -20    -    0 -    ?           ?      00:00:00 mm_percpu
1 S      0         10          2      0   80     0 -    0 -    ?           ?      00:00:00 ksoftirqd
1 I      0         11          2      0   80     0 -    0 -    ?           ?      00:00:00 rcu_sched
1 S      0         12          2      0  -40    -    0 -    ?           ?      00:00:00 migration
5 S      0         13          2      0   9     -    0 -    ?           ?      00:00:00 idle_inje
1 S      0         14          2      0   80     0 -    0 -    ?           ?      00:00:00 cpuhp/0
5 S      0         15          2      0   80     0 -    0 -    ?           ?      00:00:00 kdevtmpfs
1 I      0         16          2      0   60  -20    -    0 -    ?           ?      00:00:00 netns
1 S      0         17          2      0   80     0 -    0 -    ?           ?      00:00:00 rcu_tasks
1 S      0         18          2      0   80     0 -    0 -    ?           ?      00:00:00 rcu_tasks
1 S      0         19          2      0   80     0 -    0 -    ?           ?      00:00:00 rcu_tasks
1 S      0         20          2      0   80     0 -    0 -    ?           ?      00:00:00 kauditd
1 S      0         21          2      0   80     0 -    0 -    ?           ?      00:00:00 khungtask
1 S      0         22          2      0   80     0 -    0 -    ?           ?      00:00:00 oom_reape
```

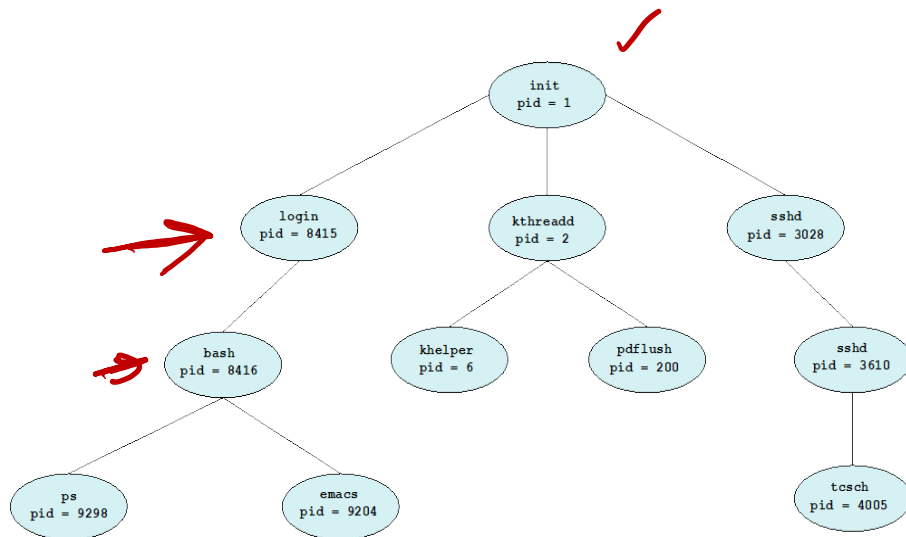
```
user1@user1-VirtualBox:~$ pstree
systemd--ModemManager--2*[{ModemManager}]
systemd--NetworkManager--2*[{NetworkManager}]
systemd--accounts-daemon--2*[{accounts-daemon}]
systemd--acpid
systemd--avahi-daemon--avahi-daemon
systemd--colord--2*[{colord}]
systemd--cron
systemd--cups-browsed--2*[{cups-browsed}]
systemd--cupsd--dbus
systemd--dbus-daemon
systemd--gdm3--gdm-session-wor--gdm-x-session--Xorg--{Xorg}
systemd--gdm3--gdm-session-wor--gdm-x-session--gnome-session-b--2*[{gnome+
systemd--gdm3--gdm-session-wor--2*[{gdm-session-wor}]
systemd--gdm3--2*[{gdm3}]
systemd--gnome-keyring-d--3*[{gnome-keyring-d}]
systemd--ibus-daemon--ibus-dconf--3*[{ibus-dconf}]
systemd--ibus-daemon--ibus-engine-sim--2*[{ibus-engine-sim}]
systemd--ibus-daemon--ibus-extension--3*[{ibus-extension-}]
systemd--ibus-daemon--ibus-ui-gtk3--3*[{ibus-ui-gtk3}]
systemd--ibus-daemon--2*[{ibus-daemon}]
systemd--ibus-x11--2*[{ibus-x11}]
systemd--2*[{kerneloops}]
systemd--networkd-dispat
systemd--packagekitd--2*[{packagekitd}]
```

pstree

Operation on processes

A process can **create** other processes to do work.

- The creator is called the parent and the new process is the child



How to share the resources?

- ① Parent & child share all resources
- ② no sharing
- ③ child can share subset of resources

How to execute?

- parent children
- ① concurrently P C
 - ② parent $\langle wait \rangle$ child
terminate

Process Creation

- Two ways to create a process
 - Build a new empty process from scratch
 - Copy an existing process and change it appropriately

✓ (+) NO wasted work

(-) setup (complex)

Process Creation

- Option 1: New process from scratch
 - Steps
 - Load specified code and data into memory; Create empty call stack
 - Create and initialize PCB (make look like context-switch)
 - Put process on ready list

Advantages:

Disadvantages:

repetition

Process Creation

- Option 2: Clone existing process and change
 - Example: fork() and exec()
 - fork() system call creates new process
 - exec() system call used after a fork() to replace the process' memory space with a new program
 - Advantages:
 - Disadvantages:

Process related system calls

- **fork()** creates a new child process
 - processes are created by forking from a parent
 - init process is ancestor of all processes
- **exec()** makes a process execute a given executable
- **exit()** terminates a process
- **wait()** causes a parent to block until child terminates

fork()

- Creates and initializes a new PCB. Creates a new address space.
- Initializes the address space with a copy of the entire contents of the address space of the parent.
- Initializes the kernel resources to point to the resources used by parent (e.g., open files)
- Places the PCB on the ready queue.
- The only difference between the child and the parent is the value returned by fork.
- Fork returns twice ✓
 - Returns the child's PID to the parent, "0" to the child

Creating a process : An example

- When you log in to a machine running Linux, you create a shell process.
- Every command you type into the shell is a child of your shell process and is an implicit fork and exec pair.
- For example, you type a command, the OS “forks” a new process and then “exec” (executes) that command.

```
user1@user1-VirtualBox:~$ pwd  
/home/user1
```

```
user1@user1-VirtualBox:~$ pwd & date
```

```
user1@user1-VirtualBox:~$ date  
Sun 24 Jan 2021 10:16:50 PM PST
```

Try!

Open the terminal and type following


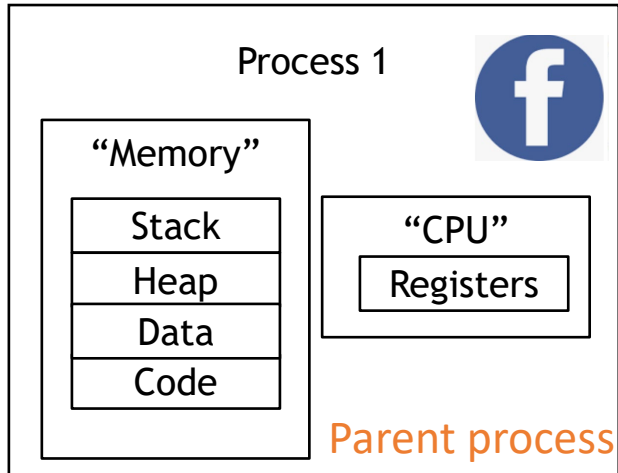
- pwd4 & date?

- pwd4 && date?

pwd && date

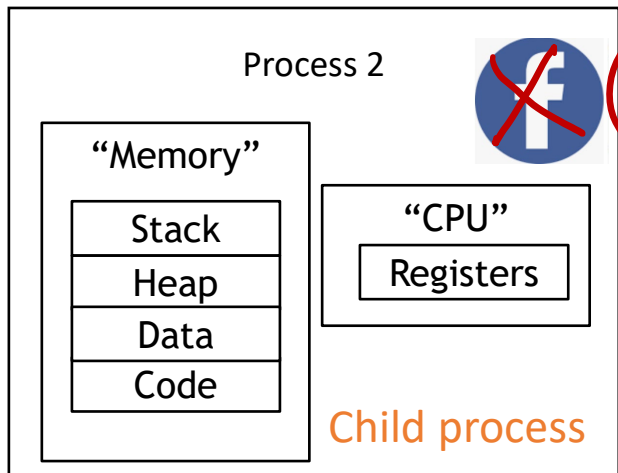
THW

fork()



fork()

exec()



- fork () creates a new child process
- Both processes run the code after the fork()
- The child is an almost exact copy of the parent except for PIDs
- fork() is called once, but returns twice
- It takes no parameters and returns an integer value.
- **Negative Value:** if creation of a child process was unsuccessful.
- **Zero:** Returned to the newly created child process.
- **Positive value:** Returned to parent process. The value contains process ID of newly created child process.

NC ↓