

CMPT 300
Operating System I
CPU Scheduling - Chapter 5

Dr. Hazra Imran

Summer 2022

Admin notes

- Quiz 3 grades available on canvas

A 2 ✓

Clicker

A non-preemptive algorithm makes a scheduling decision whenever _____.

- (1) a process changes from the blocked state to the ready state
- (2) a process requests an unavailable resource
- (3) a process terminates ✓

A. only (1)

B. (2) or (3)

C. (1) (2) or (3)

Dispatcher

- Dispatcher module gives control of the CPU to the process selected by the short-term scheduler; this involves:
 - switching context
 - switching to user mode
 - jumping to the proper location in the user program to continue executing that program
- **Dispatch latency:** time it takes for the dispatcher to stop one process and start another running

preemptive algo.

Scheduling Criteria a.k.a. Metrics

- A variety of metrics are possible.
- Many metrics involve the concept of completion

Scheduling Algorithms : Measurements

- CPU utilization: how busy are we keeping the CPU?

①

t_0 t_4 P_1 ✓
 P_2
 P_3
 P_4

- Throughput: how many processes are completed in a given unit of time

②

P_1 10 unit
30
 P_1
 P_2

- Turnaround time: how long to finish a given process?

③

- This is wall-clock time: includes waiting on I/O, kernel overhead, ...

- Waiting time: total time a process spends in ready state

④

- i.e. the process could run, but it doesn't have an available CPU

⑤

- Response time: how quickly the process begins producing output

- Scheduling algorithms can optimize for different measures

Possibilities for Optimization Criteria



- Max CPU utilization
- Max throughput
- Min turnaround time
- Min waiting time
- Min response time

First- Come First-Served (FCFS) Scheduling

- Process ready-queue is a simple FIFO
 - New processes are added to the end of the FIFO
 - Process at the front of the FIFO gets the CPU next
 - A process holds the CPU until it blocks, yields, or terminates
 - When it yields or is unblocked, it goes to the end of the FIFO
- FCFS scheduling is non-preemptive or preemptive?

non preemptive

*time sharing
system
high responsiveness*

FCFS


- No preemption of running processes
- Arrival times of processes are known
- Processors are assigned to processes on a first come first served basis
- Average waiting time can become quite large

First- Come First-Served (FCFS) Scheduling

| <u>Process</u> | <u>Burst Time</u> |
|----------------|-------------------|
| P_1 | 24 |
| P_2 | 3 |
| P_3 | 3 |

Suppose that the processes arrive in the order P_1, P_2, P_3 at time zero:

The Gantt Chart for the schedule is:



- Waiting time for each: ?
- Average waiting time: ?

$$p_1 = 0$$

$$p_2 = 24$$

$l_3 = 27$


$$(0+24+27)/3 = 17$$

First- Come First-Served (FCFS) Scheduling

Suppose that the processes arrive in the order: P_2 , P_3 , P_1 ^{3, 3, 24}

The Gantt chart for the schedule is:

17 \rightsquigarrow 3



$$P_1 = 6$$

$$P_2 = 0$$

$$P_3 = 3$$


- Waiting time for each: ?
- Average waiting time: ?

$$(6 + 0 + 3) / 3 = 3$$

convoy effect
↓
Short process behind long process


Shortest-Job-First Scheduling

- **Shortest-job-first** (SJF) scheduling orders processes based on how long their next CPU burst is expected to be
- Minimizes the average waiting time of processes
- Example: 4 processes with varying CPU-burst times:
 - 2 units, 4 units, 5 units, 7 units
- Gantt Chart of shortest-job-first ordering:



Wait times: 0, 2, 6, 11
Average wait time: 4.75

- Longest job first (for comparison):



Wait times: 0, 7, 12, 16
Average wait time: 8.75

Example of Shortest Job First (SJF)


| <u>Process</u> | <u>Burst Time</u> |
|----------------|-------------------|
| P_1 | 6 |
| P_2 | 8 |
| P_3 | 7 |
| P_4 | 3 |

- SJF scheduling chart
- Average waiting time = ?


Example of SJF

| <u>Process</u> | <u>Burst Time</u> |
|----------------|-----------------------|
| P_1 | 6 |
| P_2 | 8 |
| P_3 | 7 |
| P_4 | 3 |

- SJF scheduling chart



- Average waiting time = $(3 + 16 + 9 + 0) / 4 = 7$



Shortest Job First (SJF) Scheduling

- Biggest challenge with SJF scheduling: Predicting the length of processes' next CPU burst!
- Usually the next CPU burst length is predicted using historical data
- Common: use **exponential average** of previous bursts

Determining Length of Next CPU Burst

- Can be done by using the length of previous CPU bursts, using exponential averaging

1. t_n = actual length of n^{th} CPU burst
2. τ_{n+1} = predicted value for the next CPU burst
3. $\alpha, 0 \leq \alpha \leq 1$
4. Define: $\tau_{n+1} = \alpha t_n + (1 - \alpha) \tau_n$.

- Commonly, α set to $\frac{1}{2}$


If $\alpha = 0$, then $\tau_{n+1} = \tau_n$ and recent history has no effect (current conditions are assumed to be transient).

If $\alpha = 1$, then $\tau_{n+1} = t_n$ and only the most recent CPU burst matters (history is assumed to be old and irrelevant).

Example Burst Length Predictions

T_n


- Predicting total CPU time with $\alpha = 0.8$ and $\alpha = 0.2$.



- \ominus outliers will have erratic behavior
- \oplus very adaptive to the changes

Example Burst Length Predictions

- Predicting total CPU time with $\alpha = 0.8$ and $\alpha = 0.2$.



Clicker

What should be the value of α so that each prediction is based equally on the last observation and last prediction?

 (A) 0.5

(B) 1

(C) No idea

Shortest-Job-First Scheduling

- Shortest-job-first scheduling can be preemptive or non-preemptive
- If preemptive, called **shortest-remaining-time-first** scheduling
 - If a new job is added to the ready queue with a shorter time, it preempts the current job on the processor
- Shortest-job-first scheduling can have starvation issues
 - Some ready processes may never receive the CPU!
- Scenario: Ready queue contains short jobs and long jobs

Example of Shortest-Remaining-Time-First

- Now we add the concepts of varying arrival times and preemption to the analysis

| | <u>Process</u> | <u>Arrival Time</u> | <u>Burst Time</u> |
|--|----------------|---------------------|-------------------|
| | P_1 | <u>0</u> | 8 |
| | P_2 | 1 | 4 |
| | P_3 | 2 | 9 |
| | P_4 | 3 | 5 |

SRTS


- Preemptive SJF Gantt Chart?

Example of Shortest-Remaining-Time-First

- Now we add the concepts of varying arrival times and preemption to the analysis

| <u>Process</u> | <u>Arrival Time</u> | <u>Burst Time</u> |
|----------------|---------------------|-------------------|
| P_1 | 0 | 8 |
| P_2 | 1 | 4 |
| P_3 | 2 | 9 |
| P_4 | 3 | 5 |


- Preemptive SJF Gantt Chart*



- Average waiting time = ? msec

Example of Shortest-Remaining-Time-First


- Preemptive SJF Gantt Chart



| | Arrival Time (AT) | Burst Time (BT) | Completion Time (CT) | Turnaround Time (CT - AT) | Waiting time (TT - BT) |
|----|-------------------|-----------------|----------------------|---------------------------|------------------------|
| P1 | 0 | 8 | 17 | 17 | 9 |
| P2 | 1 | 4 | 5 | 4 | 0 |
| P3 | 2 | 9 | 26 | 24 | 15 |
| P4 | 3 | 5 | 10 | 7 | 2 |

Example of Shortest-Remaining-Time-First

- *Preemptive* SJF Gantt Chart



| | Arrival Time (AT) | Burst Time (BT) | Completion Time (CT) | Turnaround Time (CT - AT) | Waiting time (TT - BT) |
|----|-------------------|-----------------|----------------------|---------------------------|------------------------|
| P1 | 0 | 8 | 17 | 17 | 9 |
| P2 | 1 | 4 | 5 | 4 | 0 |
| P3 | 2 | 9 | 26 | 24 | 15 |
| P4 | 3 | 5 | 10 | 7 | 2 |

Average waiting time =
 $9+0+15+2 = 26/4 = 6.5 \text{ ms}$

Priority Scheduling


- Shortest-job-first is an example of **priority scheduling**
 - In SJF, the shortest job has the highest priority
- Can also assign processes fixed priorities
- Process priority is usually represented as a number *int*
 - Varies whether higher or lower numbers correspond to high priority
- Priority scheduling can be **preemptive** or **non-preemptive**

Example of Priority Scheduling

| <u>Process</u> | <u>Burst Time</u> | <u>Priority</u> |
|----------------|-------------------|-----------------|
| P_1 | 10 | 3 ✓ |
| P_2 | 1 | 1 |
| P_3 | 2 | 4 |
| P_4 | 1 | 5 |
| P_5 | 5 | 2 ✓ |

highest

- Gantt Chart ?



- Average waiting time ?


~ 8.2

Priority Scheduling

- Priority scheduling is vulnerable to starvation *aging*
 - If high-priority processes are always able to run, lower-priority ready processes will never receive the CPU
- Priority scheduling can also suffer from **priority inversion** ✓
 - Higher-priority processes are supposed to preempt lower-priority process...
 - Sometimes, in the context of resource locking, a lower-priority process can preempt a higher-priority process


Priority Inversion

- Low-priority process L starts running, and locks shared resource R .
- High-priority process H starts running, preempting L . (But L still holds resource R .)
- H needs resource R , and attempts to lock it. H blocks; L resumes.
- Medium-priority process M starts running, preempting L . M doesn't need R , and it continues to run as long as it likes.



Priority Inversion

- Because L is preempted by M , it can never finish and release R so that H can resume its execution.
- Because high-priority processes often carry out system-critical tasks, frequently has very serious consequences



Priority Inversion: Solutions

- Several solutions to priority inversion issue
- Random boosting (Microsoft Windows)
- Priority ceiling protocols
 - Every lockable resource is assigned a priority ceiling: the highest priority of any process allowed to lock it
- Priority inheritance (aka priority donation) protocols
 - If a high-priority process H is blocked waiting for a resource held by a low-priority process L , H temporarily donates its priority to L
 - A process' priority is the maximum of its own priority, and the priorities of all processes it is currently blocking

Priority Donation

- Priority donation has its own issues
- Frequently, blocked processes can form a chain

