

**CMPT 300**  
**Operating System I**  
**CPU Scheduling - Chapter 5**

**Dr. Hazra Imran**

**Summer 2022**

# Round-Robin Scheduling

5ms

15ms

- Adding time-based preemption to FCFS scheduling produces **round-robin (RR)** scheduling
  - Processes get a fixed-size **time slice** or **time quantum** on CPU
- Again, process ready-queue is a simple FIFO
  - Current process runs until it blocks, yields or terminates, or until it has used up its entire time slice.
  - When a process is moved off the CPU, it is put at end of run queue
  - Next process to receive the CPU is taken from front of the queue

System Perfo — how large time slice is.

# Example of RR with Time Quantum = 4


	<u>Process</u>	<u>Burst Time</u>
4	$P_1$	24
4	$P_2$	3
4	$P_3$	3

The Gantt chart is: ?


# Example of RR with Time Quantum = 4

<u>Process</u>	<u>Burst Time</u>
$P_1$	24
$P_2$	3
$P_3$	3

The Gantt chart is: ?



# Time Quantum and Context Switch Time



# Multilevel Queue Scheduling

- Processes can often be categorized based on their purpose and behavior, e.g.
  - System processes
  - Interactive processes
  - Interactive editing processes
  - Batch processes
- Additionally, divide processes into two main categories: foreground processes and background processes
  - **Foreground processes** need responsiveness, have small CPU bursts
  - **Background processes** have large CPU bursts, and aren't interactive

# Multilevel Queue Scheduling

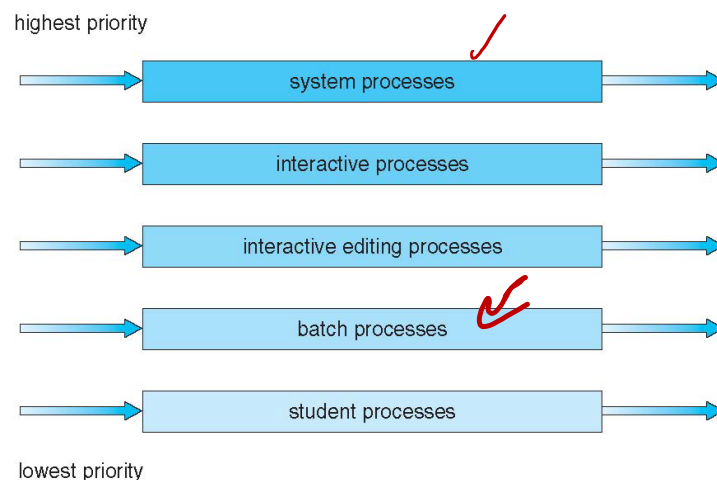
- Multilevel queue scheduling maintains a queue for each category of process
  - Queues have a decreasing priority - e.g. system processes are highest priority, batch processes are lowest priority
  - Processes are permanently assigned to a specific queue when they are started, and are not moved between different queues

# Multilevel Queue Scheduling

- Each queue has its own fixed priority
- Usually, high-priority queues always preempt low-priority
  - As long as there are system processes ready to run, they run first! ✓
  - Interactive processes only run when no system processes can run
  - etc.
  - Batch processes only run if no other processes are ready to run

Also possible to divide CPU time across subset of queues

e.g. spend 80% of CPU time running interactive processes, 20% running batch processes





# Multilevel Queue Scheduling

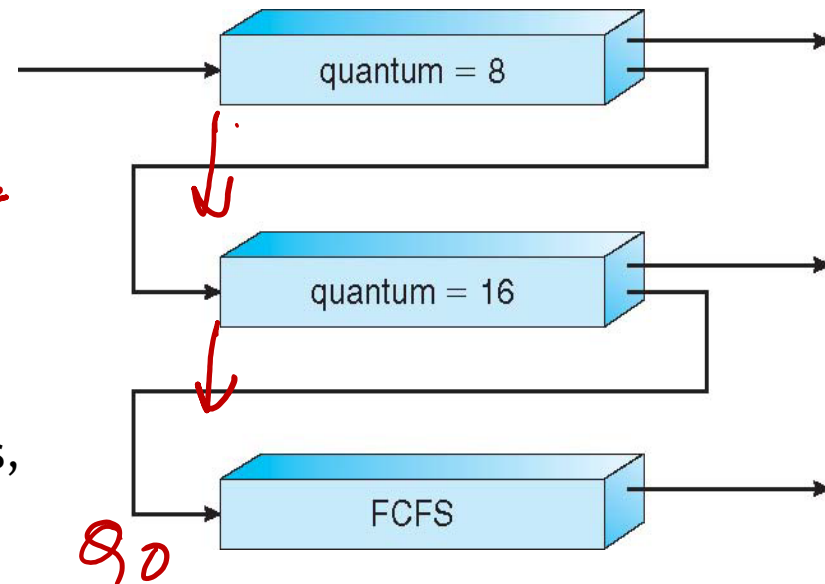
- Each queue can also have its own scheduling algorithm and parameters (e.g. time-slice size)
  - Batch processes can be run with first-come first-served scheduling, or round-robin with a very large time-slice (for runaway processes)
  - Other processes typically run with round-robin scheduling

# Multilevel Feedback Queue Scheduling

- Multilevel feedback queue scheduling allows processes to move between the different priority queues
- Multilevel-feedback-queue scheduler defined by the following parameters:
  - number of queues
  - scheduling algorithms for each queue
  - method used to determine when to upgrade a process
  - method used to determine when to demote a process
  - method used to determine which queue a process will enter when that process needs service

# Example of Multilevel Feedback Queue

- Three queues:
  - $Q_0$  - RR with time quantum 8 milliseconds
  - $Q_1$  - RR time quantum 16 milliseconds
  - $Q_2$  - FCFS
- Scheduling
  - A new job enters queue  $Q_0$  which is served FCFS
    - When it gains CPU, job receives 8 milliseconds
    - If it does not finish in 8 milliseconds, job is moved to queue  $Q_1$
  - At  $Q_1$  job is again served FCFS and receives 16 additional milliseconds
    - If it still does not complete, it is preempted and moved to queue  $Q_2$



# Multilevel Feedback Queues

- Mac OS X has multiple queues for threads, falling into four priority bands:
  - Normal (lowest priority), system high priority, kernel mode only, real-time threads (highest priority)
- Solaris uses 170 queues, divided into various categories
- Linux used a multilevel feedback queue


# Clicker

Process p is in the queue at level N-3, followed by process q at the same level. Queues at levels N through N-2 are empty. The time slice is 1 unit.

If p needs 2 units of CPU time and q needs 1 unit of CPU time, process \_\_\_\_\_ will terminate first.

(A) p

(B) q



# Clicker

When p starts executing, a new process r with a CPU time requirement of 3 units arrives at level N. The 3 processes will terminate in the order \_\_\_\_\_.

(A) r, p, q

(B) r, q, p

(C) p, q, r

$p = 2$

$q = 1$

$r = 3$

HP ✓

P, q



# Next

Ch 6 - Synchronization tool