

CMPT 300
Operating System I
Synchronization Tools- Chapter 6

Dr. Hazra Imran

Summer 2022

Synchronization Hardware

- Uniprocessors: could disable interrupts
 - Currently running code would execute without preemption
 - Generally too inefficient on multiprocessor systems
 - Operating systems using this not broadly scalable
- Modern machines provide special atomic hardware instructions
 - **Atomic** = non-interruptable
- Two types of atomic hardware instructions
 - test memory word and set value, TestAndSet()
 - swap contents of two memory words, Compare and Swap()

TAS

sort of
sequences



Test and Set Instruction



100
1000 target address

Effective behavior, but within a single instruction:

```
✓ boolean test_and_set (boolean *target)
{
    boolean rv = *target;
    *target = TRUE;
    return rv;
}
```


1. Executed atomically
2. Returns target's current value
3. Set the target's value to TRUE

Using test_and_set()

- Shared Boolean variable lock, initialized to FALSE
- Solution:

```
do {  
    while (test_and_set(&lock))  
        ; /* do nothing */  
        /* critical section */  
        lock = false;  
        /* remainder section */  
} while (true);
```

test_and_set()




```
do {  
    while (test_and_set(&lock ))  
        ; // do nothing  
    // critical section  
    lock = FALSE;  
    // remainder section  
} while (TRUE);
```

```
boolean test_and_set (boolean  
*target)  
{  
    boolean rv = *target;  
    *target = TRUE; Set the new value  
    return rv;  
}
```

Return the previous value stored in *target

Shared Boolean variable lock

- False: no process is in critical section
- True: one process is in critical section





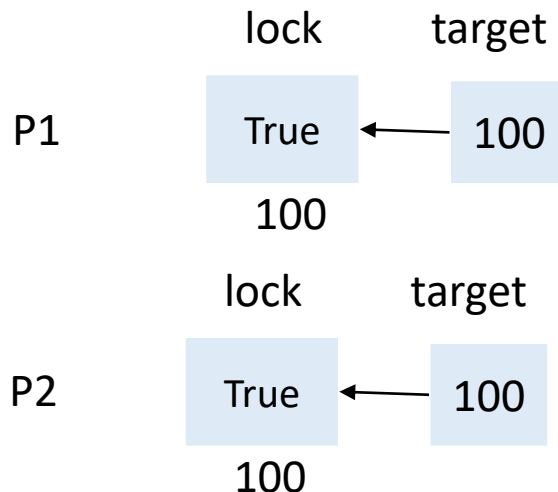
This will be set to True (new value)

Mutual exclusion using test_and_set()

```
do {  
    while (test_and_set(&lock ))  
        ; // do nothing  
    // critical section  
    lock = FALSE;  
    // remainder section  
} while (TRUE);
```

```
boolean test_and_set (boolean  
*target)  
{  
    boolean rv = *target;  
    *target = TRUE; → Set the new  
    return rv;      value  
}
```

Return the previous value
stored in *target





For P2, true will be return

test_and_set()

Does test_and_set() satisfy our Critical Section Properties?

• Mutual exclusion:	Yes	No	No Guarantee
• Progress:	Yes	No	No Guarantee
• Bounded wait:	Yes	No	No Guarantee

Compare-and-Swap

```
int compare_and_swap(int *value, int expected, int new_value) {  
    int temp = *value;  
    if (temp == expected)  Compare the actual current value with  
        the expected value compare  
        *value = new_value;  Update with the new value if the  
        current value matches expected value swap  
    return temp;  
}
```

1. Executed atomically

2. Returns the original value of passed parameter “value”

3. Set the variable “value” to the value of the passed parameter “new_value”, but only if “value” == “expected”.

That is, the swap takes place only under this condition.

Critical Sections with compare_and_swap()

- Shared integer “lock” initialized to 0;



*lock = false
0*

- Solution:

```
do {  
    while (compare_and_swap(&lock, 0, 1) != 0)  
        ; /* do nothing */  
  
    /* critical section */  
  
    lock = 0;  
    /* remainder section */  
} while (true);
```

*expected
new value*

Solution using Compare-and-Swap

```
int compare_and_swap(int *value, int expected, int new_value) {  
    int temp = *value;  
    if (temp == expected)  Compare the actual current value with  
        the expected value  
        *value = new_value;  Update with the new value if the  
        current value matches expected value  
    return temp;  
}
```

```
while(true) {  
    while(compare_and_swap(&lock, 0, 1) != 0)  
        ; /* do nothing */  
    // critical section  
    lock = 0;  
    //remainder section  
}
```

Shared variable **lock**
is initialized to 0

compare_and_swap()

Does compare_and_swap() satisfy our Critical Section Properties?

• Mutual exclusion:	Yes	No	No Guarantee
• Progress:	Yes	No	No Guarantee
• Bounded wait:	Yes	No	No Guarantee

Lock()/Unlock()

- Threads pair up calls to lock() and unlock()
 - between lock() and unlock(), the calling thread **holds** the lock
 - at most one thread can hold a lock at a time
- What happens if the calls aren't paired
 - I lock, but forgot to unlock?
- What happens if the two threads lock() different locks (lock A and B)

T1
Lock A
Lock B
shared
writes

Implementing locks

- Problem is that implementation of locks has critical sections, too!
 - the lock/unlock() must be atomic
 - atomic - executes as though it could not be interrupted
 - code that executes “all or nothing”
- Need help from the hardware
 - atomic instructions
 - test-and-set, compare-and-swap, ...

Mutex Locks (Spinlocks)

busy waiting


```
acquire() {  
    while (!available)  
        ; // busy wait  
    available = false  
}
```

```
release() {  
    available = true;  
}
```

- Both acquire and release must be done atomically
- Keeps spinning when waiting for a lock
 - May waste CPU cycles
 - No context switching occurs when thread is spinning
 - Useful especially when CS is small code
- Widely used on multiprocessor systems
 - A thread keeps spinning on one processor (waiting for lock)
 - While another thread performs CS on another processor, which will eventually release the lock for the spinning thread

Semaphores

- Semaphore S - integer variable
- Can only be accessed via two indivisible (atomic) operations
 - wait() and signal(), originally called P() and V() respectively.
 - Also known as down() and up()

```
wait(S) {  
    while (S <= 0)  
        ; // busy wait  
     S--;  
}
```

```
signal(S) {  
    S++;  
}
```

Semaphores: Usage

- **Counting semaphore**: integer value can range over an unrestricted domain
 - Can solve a wider range of synchronization problems
 - But, can still implement a Binary Semaphore
- **Binary semaphore**: integer value can range only between 0 and 1
 - Same as a mutex lock

Semaphores: Usage

Consider two concurrent processes: P1 and P2

- S1 (part of P1) must happen before S2 (part of P2)
- Semaphore “synch” is initialized to 0

P1:

```
// other code
S1;
signal(synch);
// other code
```

```
wait(S) {
    while (S <= 0)
        ; // busy wait
    S--;
}
```

P2:

```
// other code
wait(synch);
S2;
// other code
```


```
signal(S) {
    S++;
}
```

Semaphore Details

- Implementations of `wait()` and `signal()` must guarantee that the same semaphore variable is not accessed by more than one process at the same time
- `wait` and `signal` become CS (must be protected)
 - Disable interrupts (uniprocessor systems only)
 - Too expensive to disable interrupt for each core for multicores
 - Busy waiting or spinlocks (multiprocessor systems)
- With their use, we can still have the busy waiting problem. Just got shifted from application-level CS entry to semaphore's `wait` and `signal` commands, which are very short
- **Question:** So why not do the **busy waiting** in the user's critical section?
- **Ans:** User application may run for a long time, busy waiting can be costly

Semaphore Implementation with no Busy Waiting

- With each semaphore there is an associated waiting queue
- Each entry in a waiting queue has two data items:
 - value (of type integer): semaphore variable
 - pointer to a FIFO queue of processes waiting on the semaphore
- Two internal operations
 - Block** - suspend the process that invokes it (place the process in the waiting queue)
 - Wakeup** - resume the execution of blocked process (remove one of the processes from the waiting queue and place it in the ready queue)

```
typedef struct{  
    int value;  
    struct process *list;  Additional wait list  
} semaphore;
```

Semaphore Implementation with no Busy Waiting

```
wait(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this process to S->list;  
        block();  
    }  
}
```

No enough resource instances, must wait

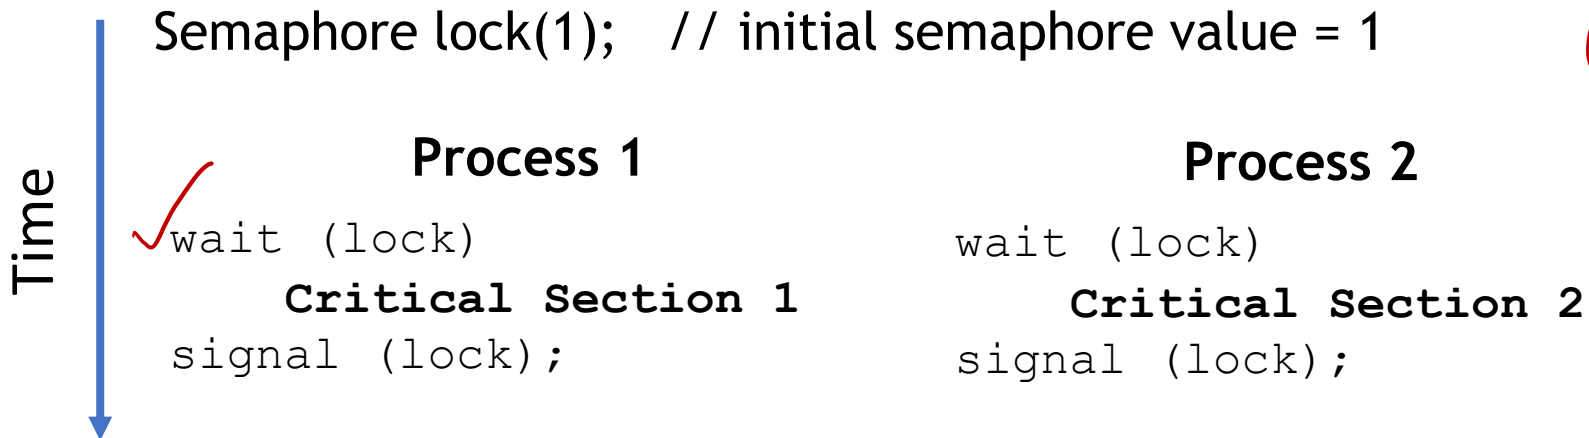

Suspend the process (Move to waiting queue)

```
signal(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a process P from S->list;  
        wakeup(P);  
    }  
}
```

Check if we can wake up a process that's been waiting

Schedule the process (Move to the ready queue)


Using Semaphores: Mutual Exclusion




- Process 1 to execute **wait** () and lock is decremented to 0
- Process 2 waits until lock becomes > 0 , which happens only when the first process executes **signal**
- Value of semaphore indicates number of waiting processes
 - lock = 0 means 1 process may be waiting
 - lock = 1 means no process is waiting

Semaphores

Let S and Q be two semaphores initialized to 1



Deadlock




Question: What could go wrong here?

Be Careful When Using Semaphores

- Incorrect order accessing multiple semaphores across processes
 - May cause deadlock
- signal (mutex) wait (mutex)
 - Multiple processes can access CS at the same time
- wait (mutex) ... wait (mutex)
 - Processes may block for ever
- Forgetting wait (mutex) or signal (mutex)
 - Various problems, inconsistent data, ...

Summary

- Synchronization: coordinate access to shared data
- Race condition 
 - Multiple processes/threads manipulating shared data and result depends on execution order
- Critical section problem
 - Three requirements: mutual exclusion, progress, bounded waiting
 - Software: Peterson's Algorithm
 - Hardware: Test_and_set, Compare-and-Swap and others
 - Busy waiting (or spinlocks)
 - Mutexes, semaphores

