

CMPT 300
Operating System I
Memory Management - Chapter 9

Dr. Hazra Imran

Summer 2022

Outline

- Introduction
- Virtualization
- **Registers**
- Segmentation
- Swapping
- Fragmentation
- Free Space Management
- Paging


Base and Limit Registers

virt addr — phy address

A **base** and a **limit** register define the logical address space

- Physical address = **Base + Virtual Address**
 - Base: starting address in memory (physical)
 - Virtual address: offset into the address space

*① specialized
② super fast hw components*

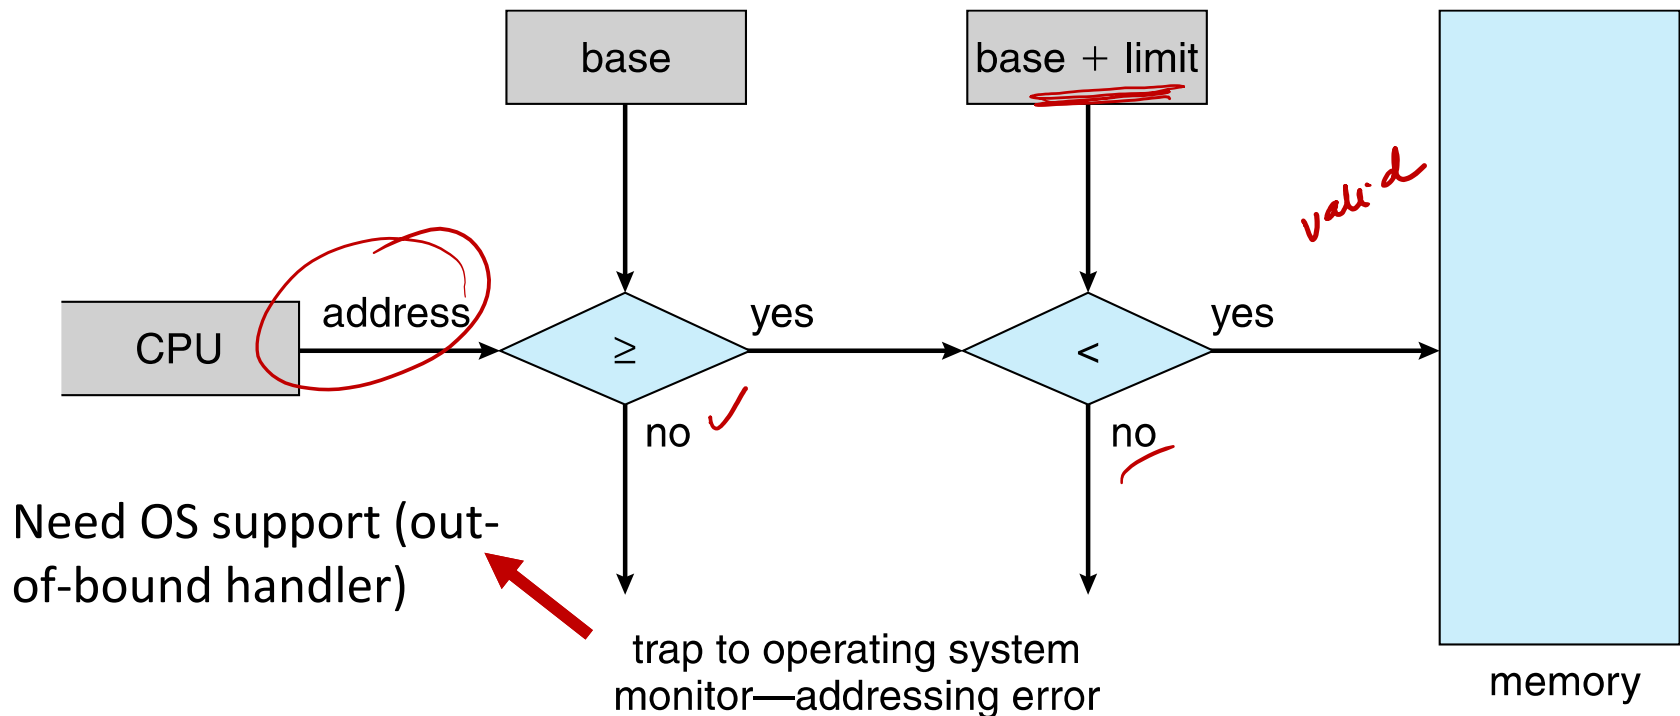


- Supported by memory management unit (MMU)
- Privileged instructions for modifying these registers

Base and Limit Registers

A **base** and a **limit** register define the logical address space

- Physical address = **Base + Virtual Address**
 - Base: starting address in memory (physical)
 - Virtual address: offset into the address space



Base and Limit Registers

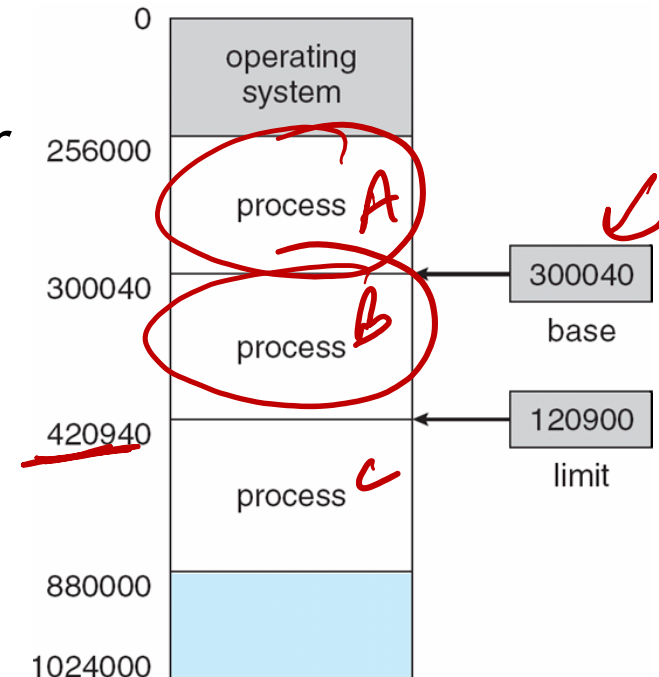
A **base** and a **limit** register define the logical address space

- Physical address = Base + Virtual Address

Limitations:

- Logical address space cannot be larger than physical address space
 - I.e., entire process fits in memory
- Logical address space must be stored **contiguously**
- May waste space
 - Space between heap and stack

Better alternative: segmentation




Outline

- Introduction
- Virtualization
- Registers
- **Segmentation**
- Swapping
- Fragmentation
- Free Space Management
- Paging

Segmentation: Generalized Base+Limit

Programmer's view of memory:


- A program is a collection of segments
- A segment is a logical unit such as:
 - Main program ✓
 - Functions, procedures
 - Method
 - Object
 - Local variables, global variables
 - Stack
 - Symbol tables ✓
 - Arrays ✓




Segmentation: Have a base + limit pair per logical segment

Segmentation: Generalized Base+Limit

Have one base + limit pair per logical segment




user space



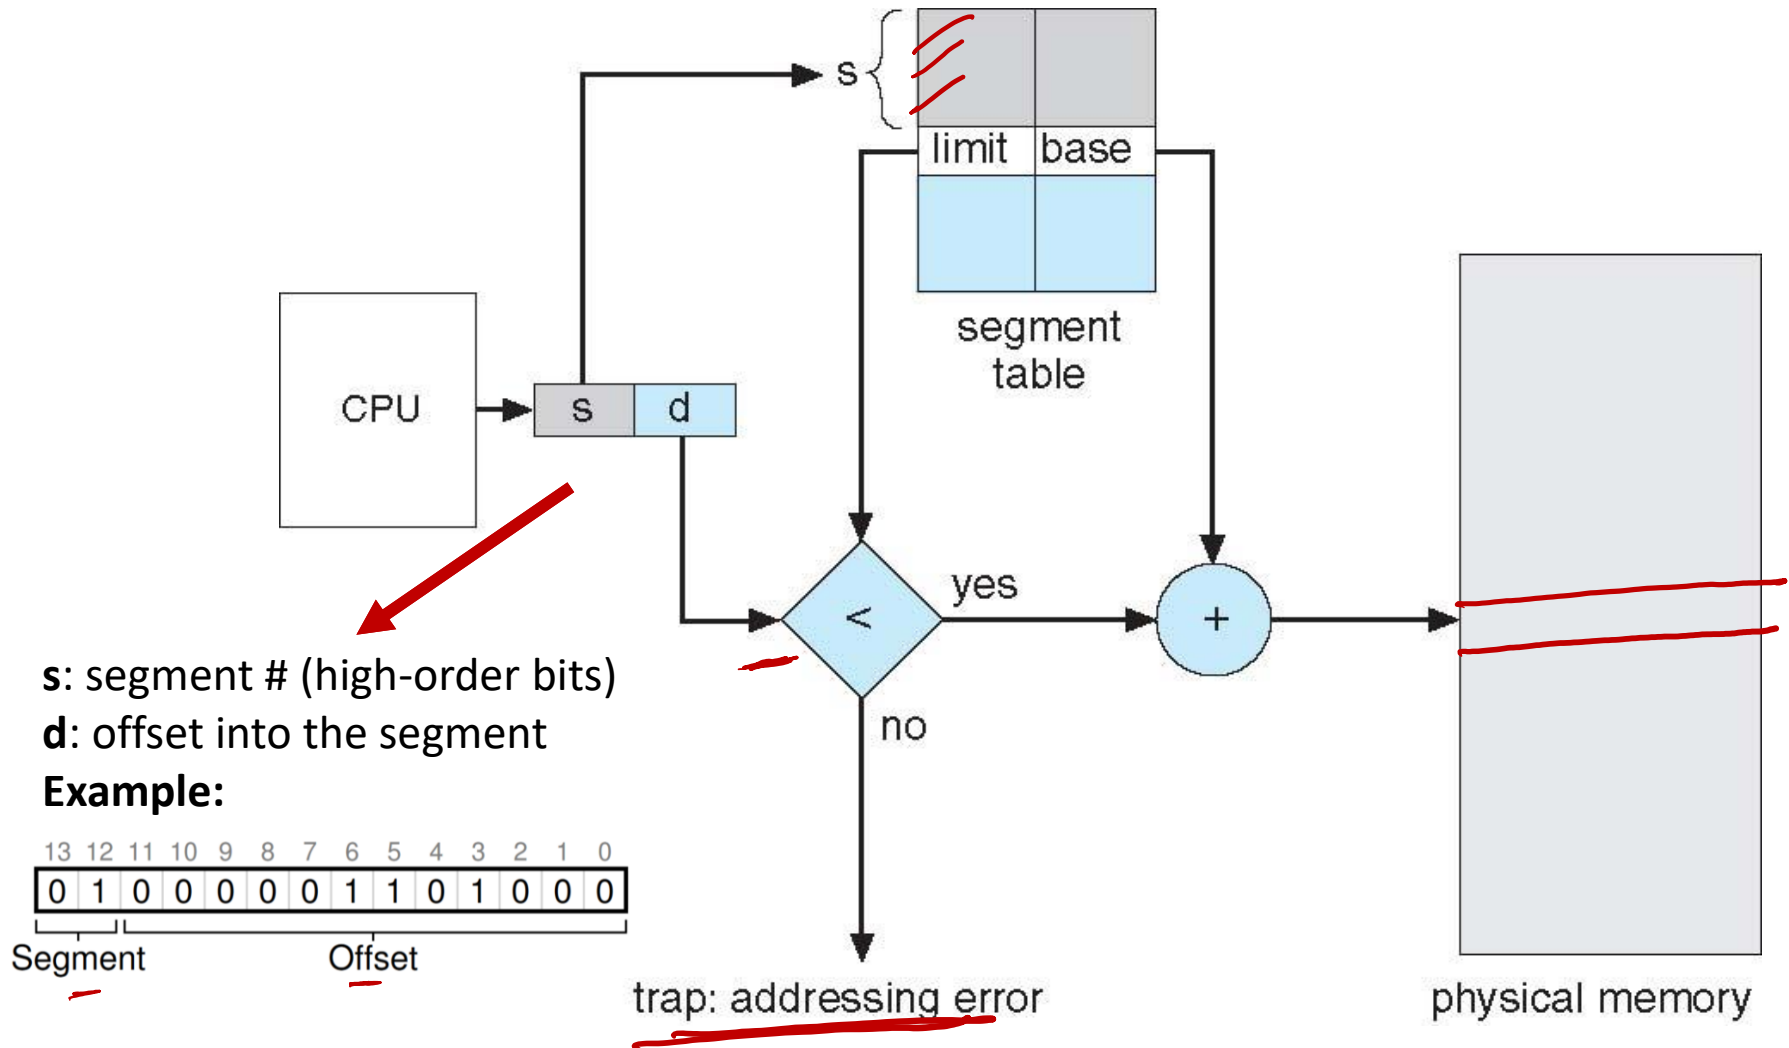
physical memory space

Segmentation Architecture


- Logical address has the form: <segment-number, offset>
- Segment table
 - Maps logical address to physical address
 - Each entry has
 - Base: starting physical address of segment
 - Limit: length of segment *(offset)*



Segmentation Hardware




Segmentation Example



	limit	base
0	1000	1400
1	400	6300
2	400	4300
3	1100	3200
4	1000	4700

segment table



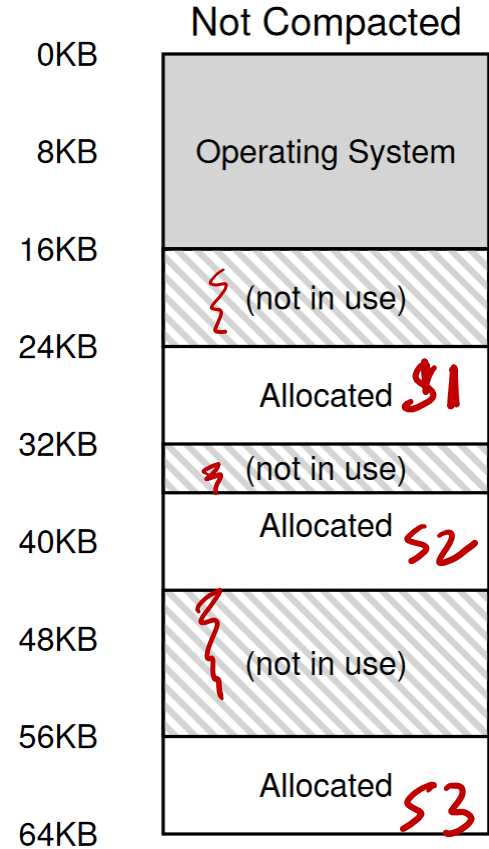
Issues in Segmentation

Issue 1: Need to save and restore segment registers upon context switch


Issue 2: Managing free spaces:

- Segments vary in length, OS needs to find free space in memory for new address space
- This can lead to **external fragmentation**
 - Physical memory becomes full of small holes of free spaces
 - hard to allocate new segments or grow existing ones
- **Solutions:** free space management schemes

Fragmentation example:




Summary

- Address space
 - Easy-to-use abstraction of memory
 - Gives processes the illusion that it owns the entire system
 - Provides isolation and protection between processes
- Basic approach: base and limit registers
 - Base: starting physical address
 - Limit: size of segment
 - Need privileged instructions to modify base and limit registers
 - Various limitations
-  Segmentation: generalized base+limit registers approach
 - One pair of base+limit registers per segment
 - Main issue: external fragmentation

Segmentation

- The process address space is a set of disjoint pieces (text; data; stack; heap), i.e. no overlapping of the pieces
- Some pieces dynamically grow and/or shrink (stack; heap; text if loaded on demand)
- The programmer does not care how pieces are organized in the address space
- These pieces are **Segments**



Segmentation

- The logical address space is a collection of segments. A logical address is decomposed as:
 - A **segment number**
 - An **offset in the segment**
- The compiler/language interpreter handles the segments and the logical addresses are built appropriately
- Typical segments used by a C compiler
 - Text, data, heap ,stacks. standard C library

Segmentation

- Logical address = Segment number + Offset
- Protection mechanisms need to be set up.
- A **segment table** with one entry per segment number
 - Base: Starting address of the segment
 - Limit: Length of the segment
- The segment table is stored in memory
 - ① • A Segment-Table Base Register (STBR): Points to the segment table address
 - ② • A Segment-Table Length Register (STLR): Gives the length of the segment table; Makes it easy to detect an invalid segment offset

MMU and Segmentation

- Segmentation is easy
- Reserve bits (e.g., the left-most ones) in the logical address to reference a segment (the segment bits)
- Lookup the segment table to find out the segment's base/limit value segmentation

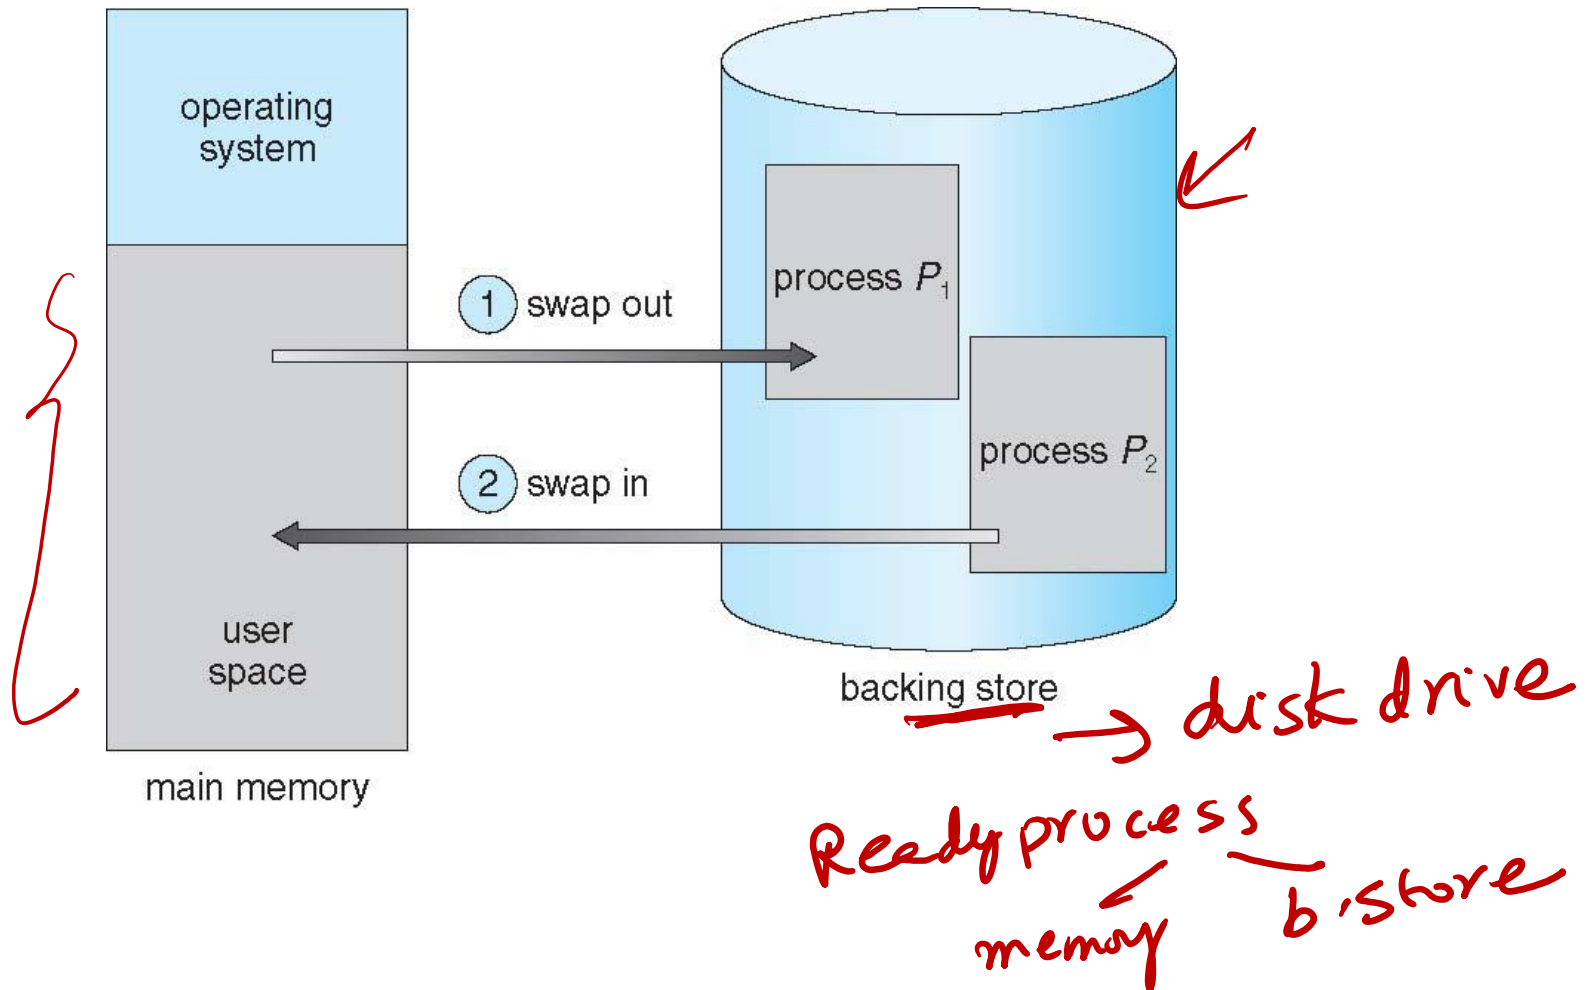
Demand on memory space

- Considering that multiprogramming enables our computer to simultaneously serve many processes, and memory size is much smaller than disk drive.
- Then what if jointly all processes ask for more memory space than is available?

Outline

- Introduction
- Virtualization
- Registers
- Segmentation
- **Swapping**
- Fragmentation
 - Free Space Management
- Paging

Schematic View of Swapping



Swapping

- Does the swapped-out process need to swap back into same physical addresses?
- ✓
- Depends on address binding method
 - Plus consider pending I/O to / from process memory space
- Modified versions of swapping are found on many systems (i.e., UNIX, Linux, and Windows)
 - Swapping normally disabled
 - Started if more than threshold amount of memory allocated
 - Disabled again once memory demand reduced below threshold

Swapping

- What if we want to load another process that would not fit in memory?
- We must save the address space of one (or more) processes from RAM to a “backing store” (the disk)
- Moving processes back and forth between main memory and the disk is called swapping
- When a process is swapped back in, it may be put into the same physical memory space or not
- (But this is no problem....thanks to address virtualization)
- With swapping a process can “be in RAM” or “be on Disk”.
Therefore, a context-switch can involve the disk!!



Swapping and DMA

- With swapping, at any time the OS could kick a process out of RAM and save it to disk
 - This raises a concern with Direct Memory Access (DMA)
- Consider a process that has initiated a DMA operation and is swapped to disk
 - The DMA controller may have no idea and happily continues to write data (into some other process' address space, which has replaced that of the one that was swapped out!)
 - Operating systems must deal with this (because DMA is so useful we can't leave without it)
 - One option could be: never swap a process engaged in DMA
 - In fact, OSes do something else ("paging", next topic)

Context Switch Time including Swapping

- If next processes to be put on CPU is not in memory, need to swap out a process and swap in target process
- Context switch time can then be very high
- 100MB process swapping to hard disk with transfer rate of 50MB/sec
 - Swap out time of 2000 ms
 - Plus swap in of same sized process
 - Total context switch swapping component time of 4000ms (4 seconds)
- Can reduce if reduce size of memory swapped - by knowing how much memory really being used
 - System calls to inform OS of memory use via `request_memory()` and `release_memory()`

Context Switch Time and Swapping (Cont.)

- Other constraints as well on swapping
 - Pending I/O - can't swap out as I/O would occur to wrong process
 - Or always transfer I/O to kernel space, then to I/O device
 - Known as double buffering, adds overhead
- Standard swapping not used in modern operating systems
 - But the modified version is common
 - Swap only when free memory is extremely low

The Bad News about Swapping

- The disk is sloooooooooow (even if it's an SSD)
- Several ways to cope with slow disks have been used:
 - An OS could swap in/out only processes with small address space (rather than processes with large address space)
- One can dedicate a disk/partition to swapping (so as to minimize disk seeks on a hard drive)

The Bad News about Swapping

- One approach is to just **not swap**.
- Swapping should be an exceptional occurrence
 - In older OSes swapping was user-directed (e.g., Windows 3.1)
 - Swapping is now often disabled (e.g., on laptops)
- If the normal mode of operation of the system requires frequent swapping, the system is in trouble (buy more RAM!)
 - But perhaps it's just a temporary rare load spike?
- A key solution is to not swap whole address spaces (“paging”)

Where are we?

We now have a whole bunch of **mechanisms**:

- We know how to allocate memory “portion” to each process
- We know how to reduce address spaces
- We know how to swap processes in and out of memory

We now need a policy to decide how to place each portion in memory:


- We want to have as many process address spaces in memory as possible
- We want to minimize swapping
 - What is a good policy?

Outline

- Introduction
- Virtualization
- Registers
- Segmentation
- Swapping
- **Fragmentation**
 - Free Space Management
- Paging

Contiguous Memory Allocation

- Segments vary in length; OS needs to find free space in memory for new address space
- As processes/segments are created and destroyed, memory becomes full of holes of free spaces
- Holes have different sizes and scattered in the memory



Contiguous Allocation Policies ✓

Which hole to choose for a process?

- OS maintains information (lists) on allocated space and free space (holes)

First-fit: Allocate first hole that is big enough

Best-fit: Allocate smallest hole that is big enough

- Must search entire list, unless ordered by size
- Produces the smallest leftover hole

Worst-fit: Allocate the largest hole

- Must also search entire list
- Produces the largest leftover hole



Pros and cons of contiguous allocation:

- Pros: *simple to implement*
- Cons: *→ memory fragmentation*

*↓
NC*