

**CMPT 300**  
**Operating System I**  
**Memory Management - Chapter 9**

**Dr. Hazra Imran**

**Summer 2022**

# Outline

- Introduction
- Virtualization
- Registers
- Segmentation
- Swapping
- Fragmentation
  - Free Space Management
- Paging

# Logical address vs Physical address

- Logical address - generated by the CPU <sup>virtual address</sup>
- Physical address - address seen by the MMU
- The term Logical Address Space is used for the set of all logical addresses generated by a program's perspective.

CAS & PAS same in compile-time  
and load time  
address binding  
schemes

CAS & PAS differ in  
execution-time  
based schemes

# Recap

## Main-Memory Management

- MM is a large array of words or bytes.
- Each word or byte has its own address.
- MM is a repo of quickly accessible data shared by the CPU and I/O devices.

The major activities of an operating system regarding MM are:


- Keep track of which part of memory is currently being used and by whom.
- Decide which processes are loaded into memory when memory space becomes available.
- Allocate and deallocate memory space as needed.

# External Fragmentation

Total memory space enough for allocation, but it is not contiguous

## Solution - Compaction:

- Copy data to one contiguous chunk of memory
- Main disadvantage of compaction: slow



# Internal Fragmentation

- Allocated memory may be larger than requested
  - E.g., asked for 5 bytes (through malloc), allocated 8 bytes



3 extra bytes left  
unused

- Occurs when the system has fixed allocation sizes
- The hole is *internal* to the allocated memory
  - Compaction cannot help: OS/allocator has no control once the memory is handled to the requesting process

# Memory Allocation Mechanisms

- Techniques used by contiguous allocation policies
  - Best-fit, worst-fit, first-fit...
- Useful for both OS and user-space memory allocators
  - E.g., malloc (malloc + free interfaces)
- Mechanisms
  - Splitting
  - Coalescing
  - Tracking allocated memory
  - Tracking free memory

# Splitting

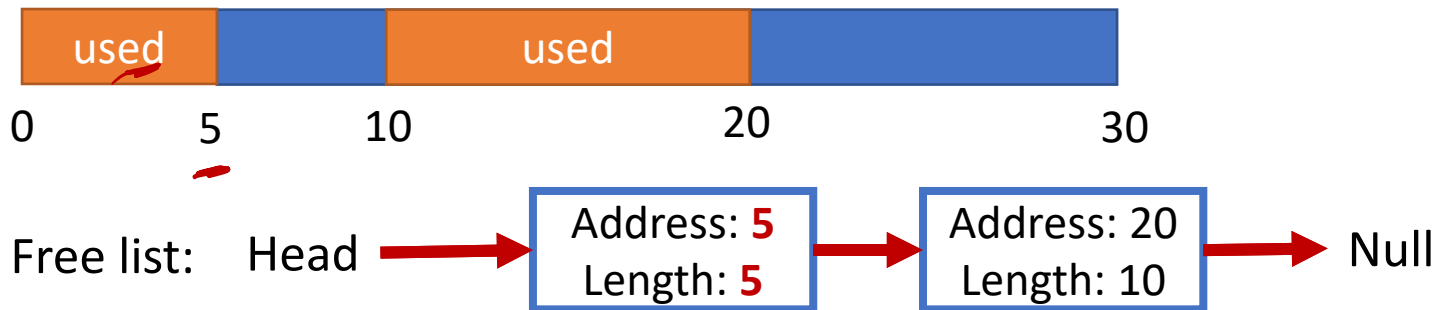
Find a free chunk of memory, split it into two

- Useful when requested memory size is smaller than the size of free memory

1. Initially: 30 bytes space in total, 10 bytes already allocated



2. After splitting the first 10 bytes for a new allocation of 5 bytes:



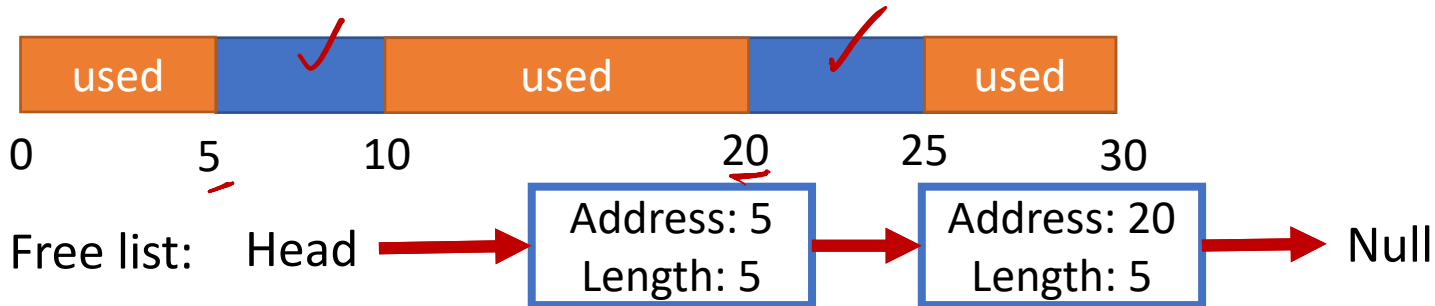


# Coalescing

Merge returned (adjacent) memory chunks into a big chunk

- Might be required for bigger allocations


1. Initially, 3 chunks being used (allocated)



2. First and third chunk freed:




3. After coalescing adjacent free chunks:




These two can also be merged

# Tracking Allocated Memory

## Memory allocator interfaces:

- Allocation: `void *allocate(int size);`  No “size” parameter
  - E.g., malloc
- Deallocation: `void deallocate(void *memory);`
  - E.g., free
  - Size information needed for managing free space (“holes”)
  - Solution: store extra information (e.g., size) in a header block together with the allocated memory

Example: `ptr = malloc(10);`



Returned the caller: `ptr`

Used by allocator: `hptr = ptr - sizeof(Header)`


- Header size included in free list
- Not visible to user

# Tracking Free Memory

- Free list needed for tracking free memory chunks
  - Require memory to store the nodes, too
  - But there is no allocator to use - we are building one!
  - Solution: embed the list in memory


## Example:

- Total memory (heap) size: 256 bytes
- Initially no allocation
  - Conceptually, the “free list” has only one entry
  - Head pointer points to the beginning of the memory space
- Use split/coalescing techniques for allocation/deallocation




# Tracking Free Memory

Initial state:

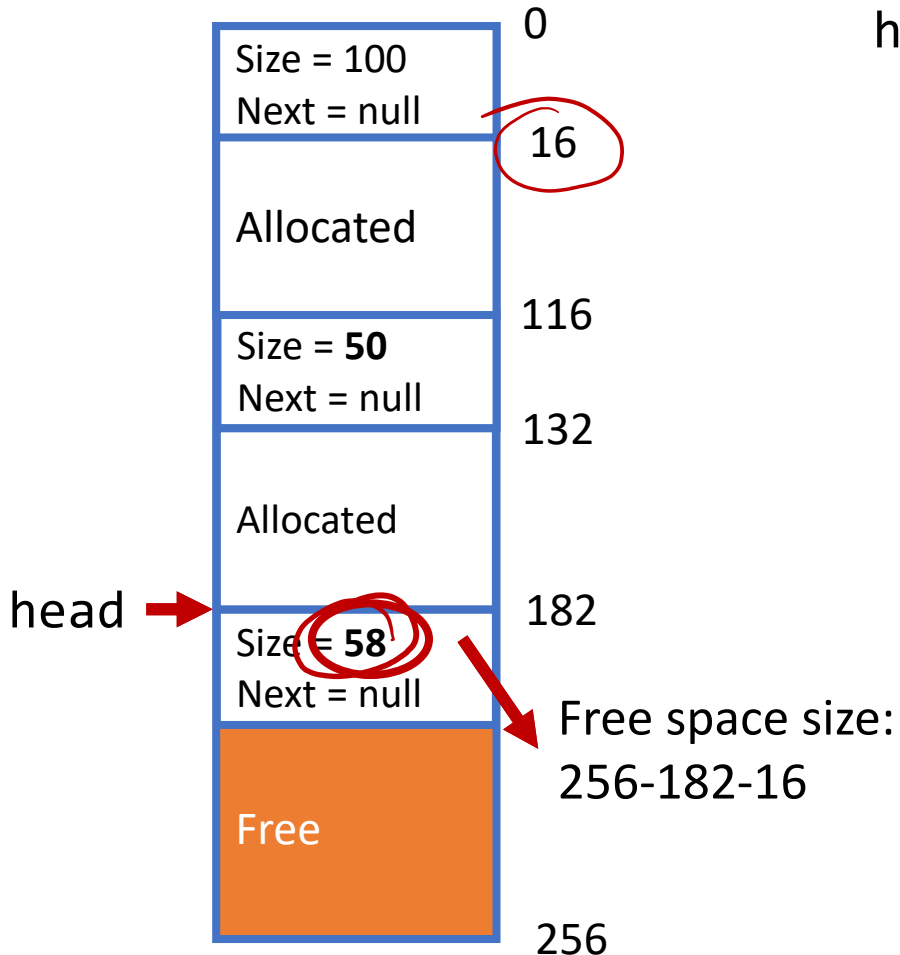


After allocating 100 bytes:

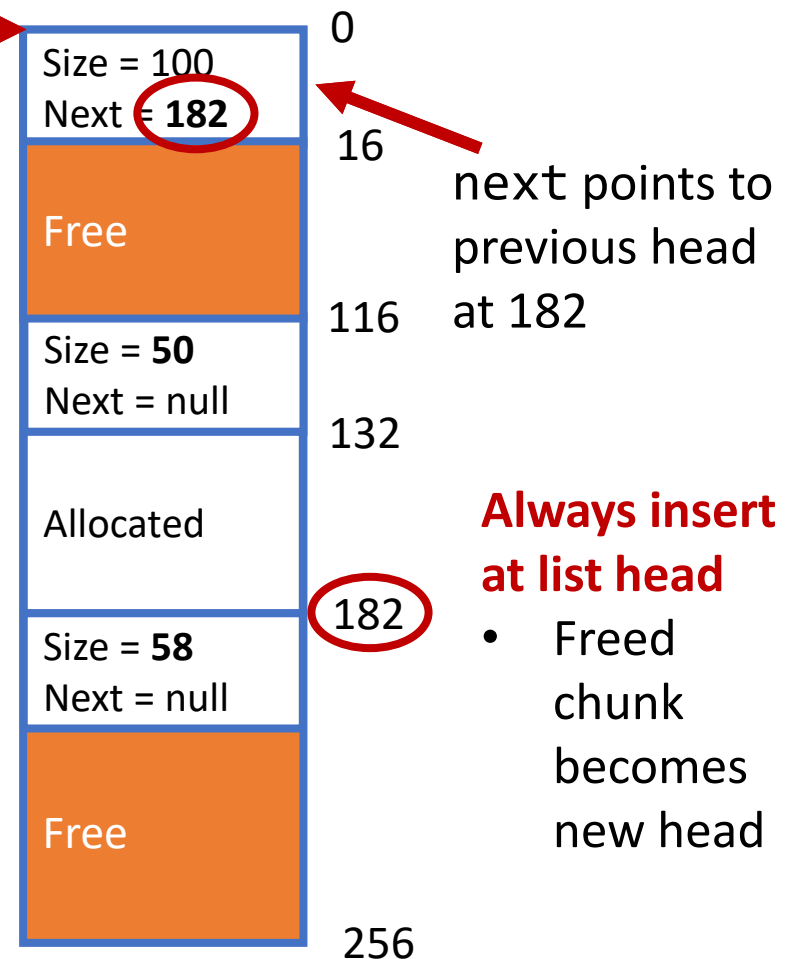


# Tracking Free Memory

Allocating another 50 bytes:




Freeing the first allocated chunk (100 bytes)  
head →



# Tracking Free Memory


Freeing the next 50 bytes:



**Q: Allocate 150 bytes?**

A: No way to allocate –  
space fragmented in  
three pieces.  
Coalescing needed.

Coalescing:




# Let's summarized

- Objective: Allocate a contiguous slab of memory to each process so that they can store their address space
- The mechanisms are “easy”
- MMU for memory protection and relocatable code Segmentation
- Dynamic Memory / Dynamic Loading / Dynamic Linking for smaller address spaces
- ...
- But finding a good policy is really hard
- For process picking, hole picking, placement in hole
- Fragmentation is unavoidable and wastes RAM
- **Bottom Line:** Contiguous memory allocation is difficult

# Contiguous vs. Non-contiguous Allocation

## Contiguous allocation


- Addresses for a segment or process must be contiguous



Segmentation

## Non-contiguous allocation

- Paging: underlying physical memory and segment decoupled



Segmentation with paging



# Outline

- Introduction
- Virtualization
- Registers
- Segmentation
- Swapping
- Fragmentation
  - Free Space Management
- **Paging**

# Paging

PM - frames  
LM - pages

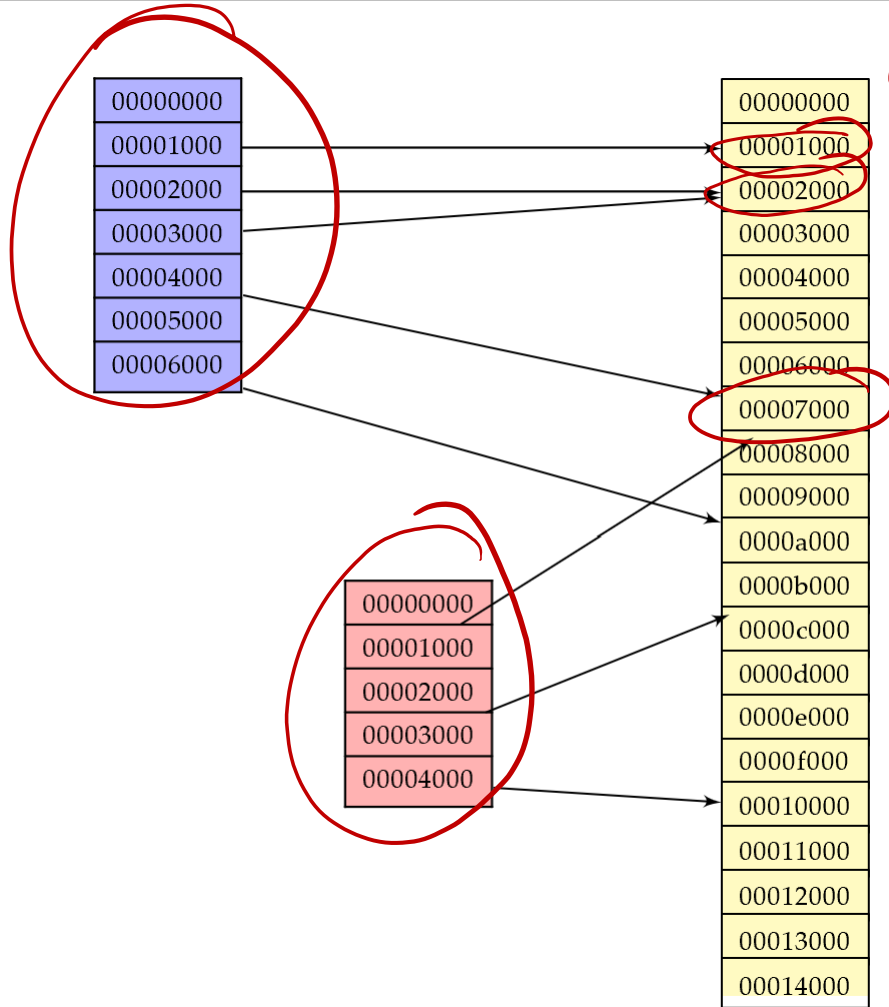
- Process is allocated memory wherever it is available
- Divide physical memory into fixed-sized blocks called **frames**
  - Size is the power of 2 (typically 512 bytes to few GBs)
  - OS keeps track of all free frames
- Divide logical memory into blocks of the same size called **pages**
- A page can be put at any frame
- For a program of size  $n$  pages, need to find  $n$  free frames
- Use page tables to translate logical to physical addresses

# Paging

External frag — NO  
Inter frag — YES

- Permits physical address space of a process to be physically non-contiguous
  - The process is allocated physical memory whenever the physical memory is available - no external fragmentation
    - What about Internal fragmentation?
  - This avoids compaction.
- program / — P3  
      P1 P2
- Physical blocks - Frame
  - Logical blocks - Pages ✓
  - Size of frames and pages is defined by hardware (power of 2 to ease calculations)
  - An address is determined by:
  - page number (index into table ) + offset
    - <-- mapped into --> base address (from table ) + offset.

# Paging



Yellow is the physical RAM pages (called *frames*).

Blue is process 1's virtual page table.

Red is process 2's virtual page table.

Each process  
↓  
PT

# Lets take an example

offset - house

street - Page #



Assume:

$\langle 2, 4000 \rangle$

Virtual address is 32 bits, page offset is 12 bits, physical address is 30 bits



Logical address

Physical address



Logical Page number	Physical Frame #
0	2
1	0
2	14
...	
1048575	

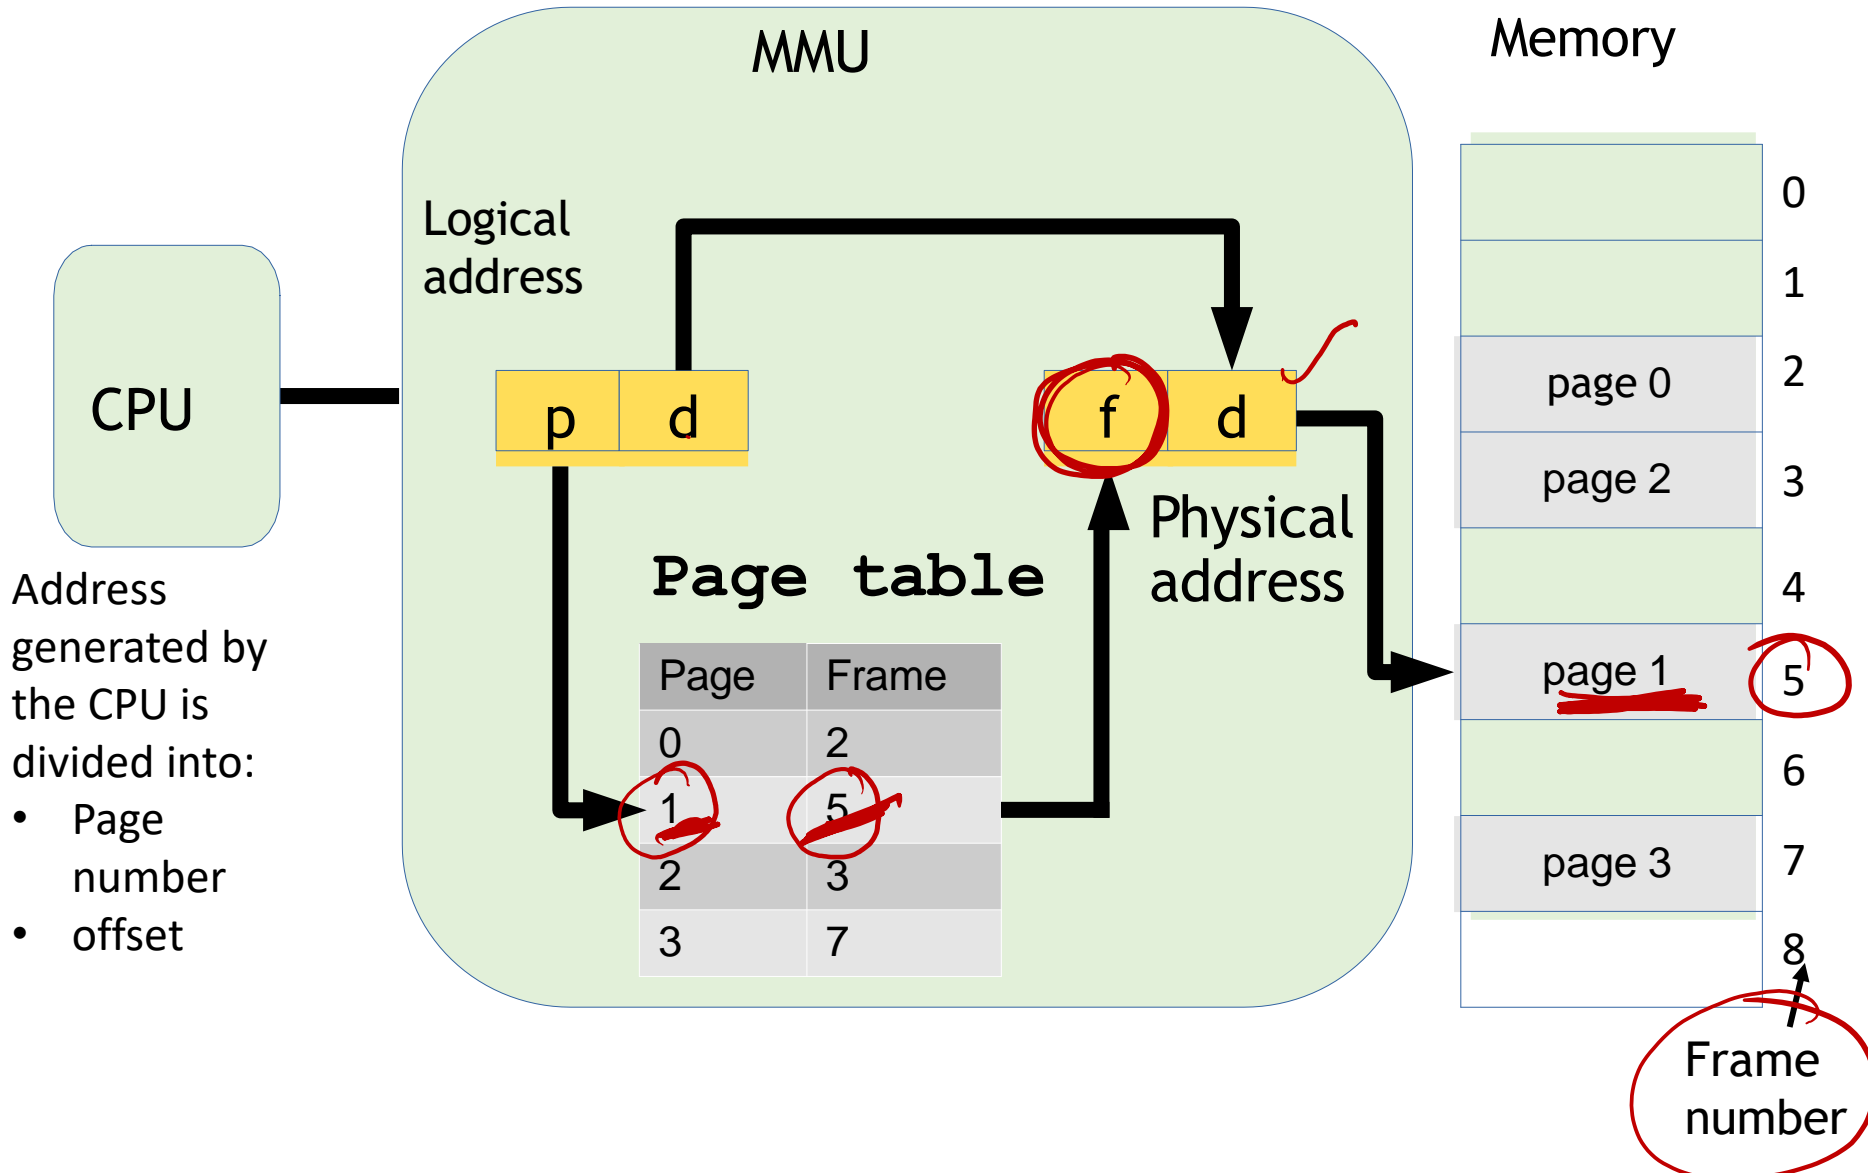
$2^{20}$  entries



Page number	
0	
...	
14	0 1 2 ... 4000 ... 4095
...	
262,143	


$2^{18}$  entries

# Page address translation



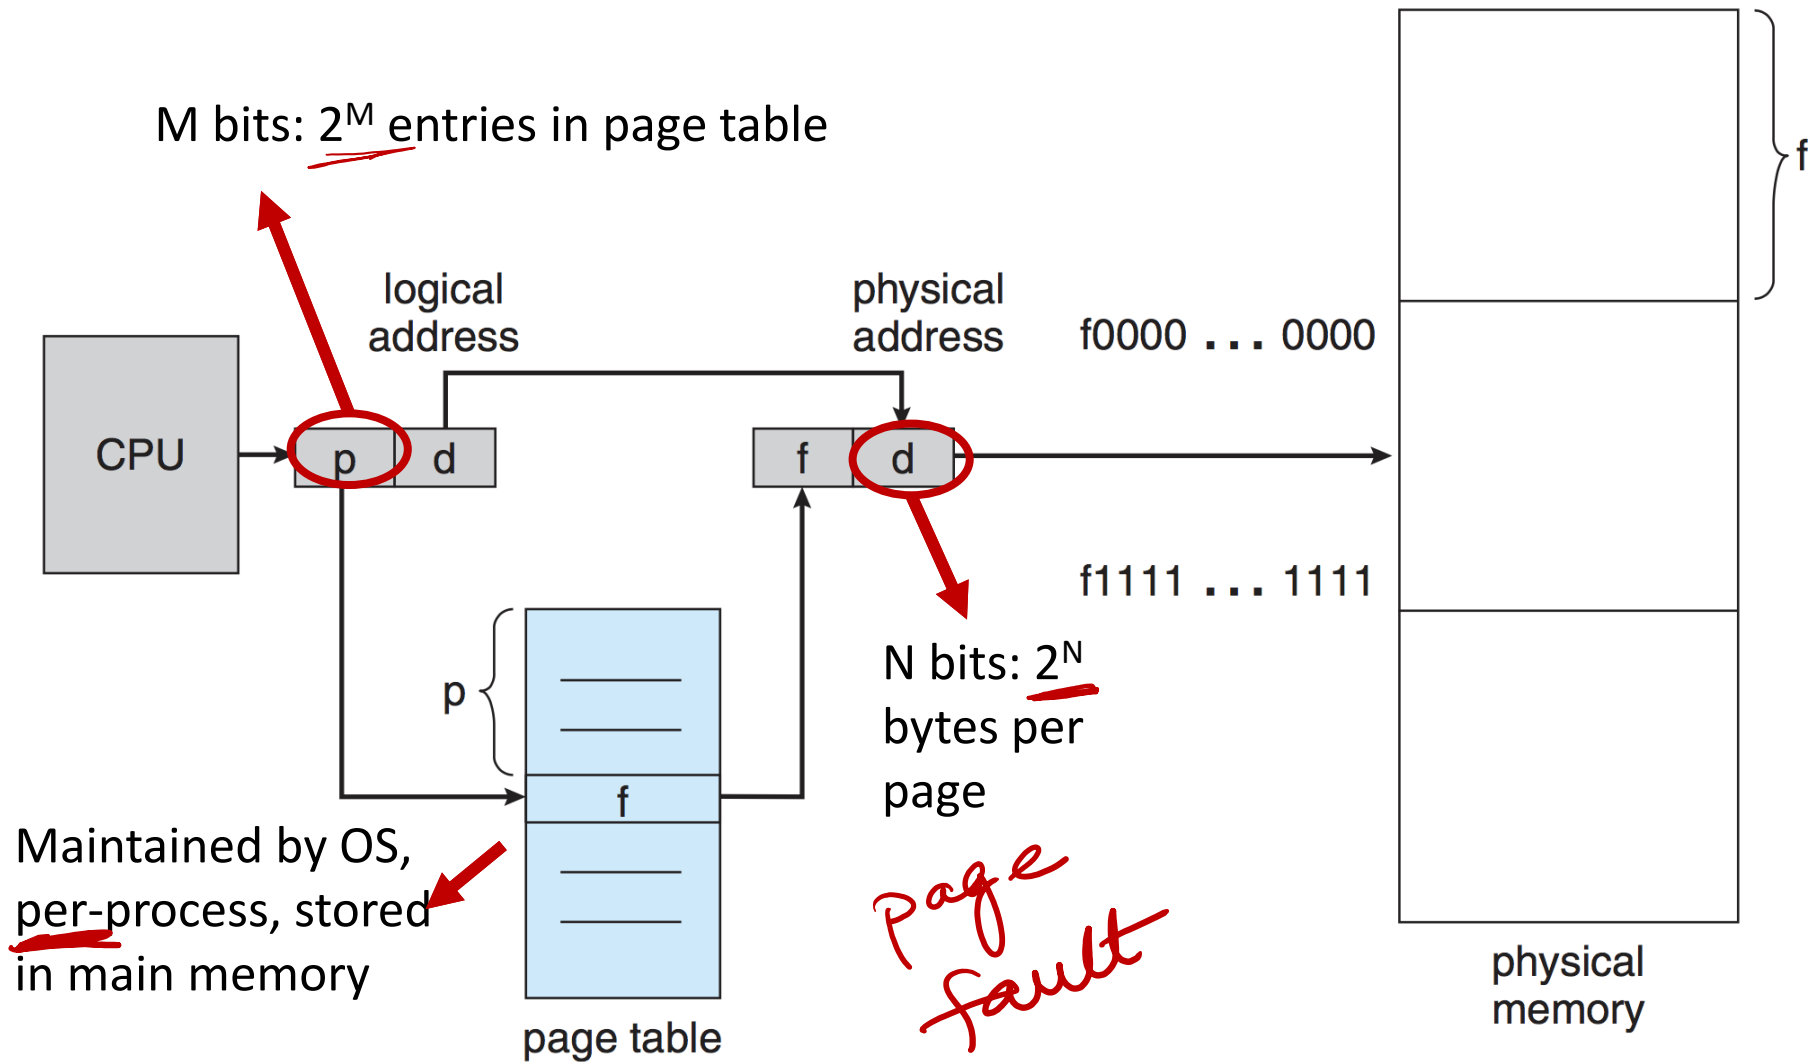
# Virtual Addresses

- Address generated by CPU is divided into:
  - Page number (p): used as an index into a page table, which is an array of entries for translating page # to frame #
  - Page offset (d): combined with frame # addr to define the physical address that is sent to the memory unit



- Logical address space =  $2^m$  bytes
- Page size =  $2^n$  bytes
- Number of pages =  $2^{m-n}$

# Address Translation in Paging





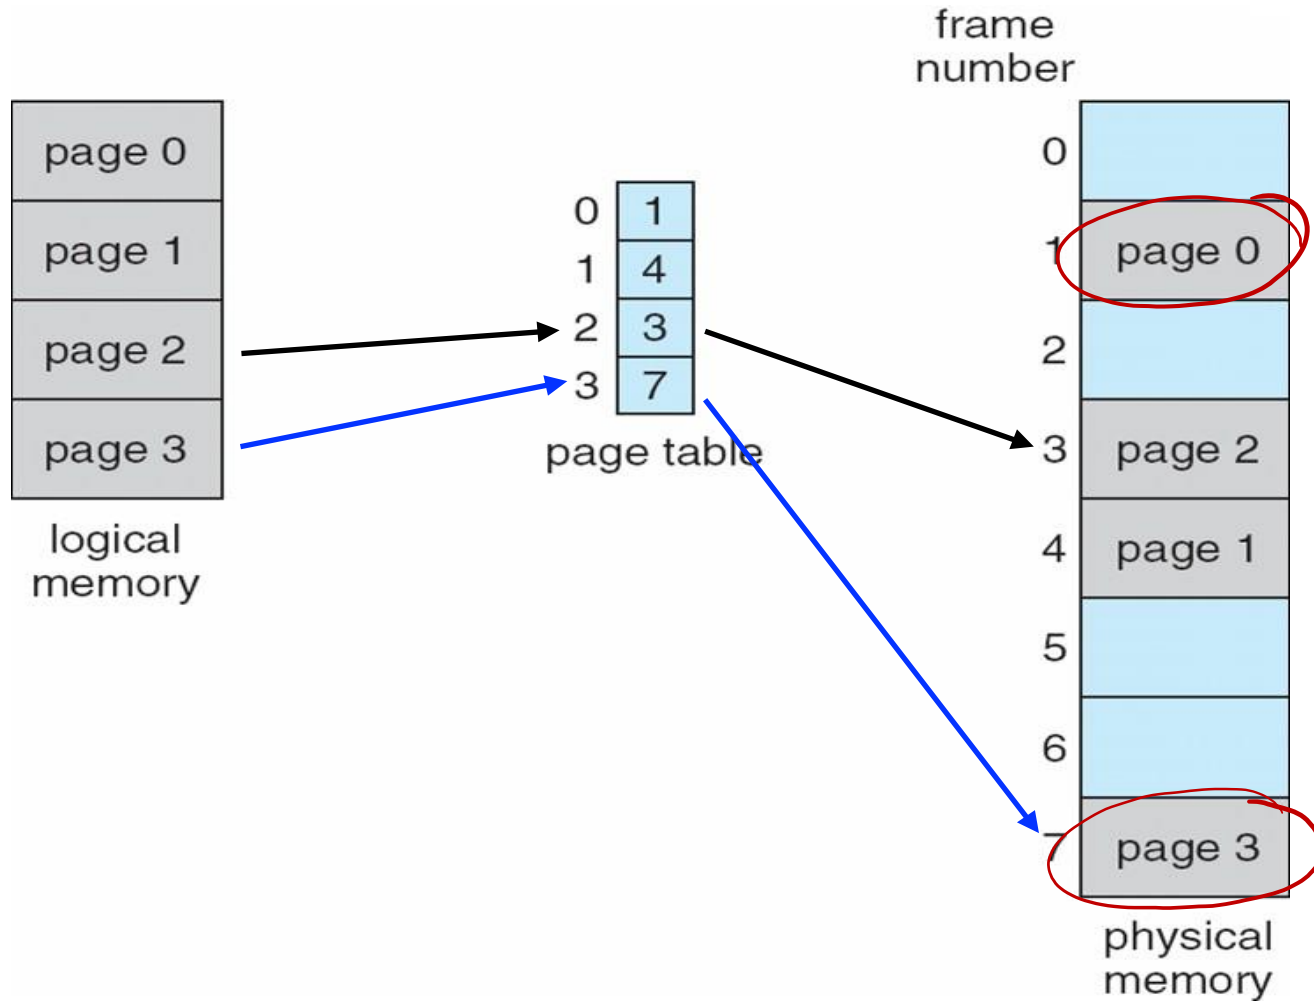
# Try!

Given 4KB page size and a 32-bit virtual address, calculate


- m (# of bits for representing virtual address space) 32
- n (# of bits for offset within a page) 12
- the number of bits for representing page #  $(32-12)$  20

$$4KB = 4096 \text{ bytes}$$
$$4KB = 2^{12}$$
$$4KB = 4 \times 1024$$

# Paging Example



# Paging Example



Ques like:

Virtual address space: 16 bytes ( $m = 4$ )

Memory size per page: 4 bytes ( $n = 2$ )

- (1) Infer the page table and
- (2) use it to translate the virtual address

page number	page offset
$p$	$d$
$m - n$	$n$

Setup like

- Page size: 4-byte ( $2^2$ )
- Physical memory: 32-byte ( $2^5$ )
- Logical memory: 16-byte ( $2^4$ )

# Address Translation with Paging

- For each process, there is an address  $A$

$$\text{page number} = A / \text{page\_size}$$

$$\text{offset} = A \bmod \text{page\_size}$$

- This is the page number within the process address space
- e.g. address in the process,  $A = 10,000$  and page size =  $4k$
- page number =  $10000 / 4k = 10,000 / 4096 = 2.\text{xxx} =$  truncate to 2
- this is the distance from the beginning of the page
- page offset =  $10000 \bmod 4k = 10,000 \bmod 4096 = 1908$

↑ NC  
1